


# Warmup #1

Bill Cheng

*<http://merlot.usc.edu/william/usc/>*




# Discussion Sections



## **IMPORTANT:**

- please understand that discussion section material are **NOT** substitute for reading the specs and the grading guidelines
  - you are expect to read the **specs**
  - you are expect to read the **requirements** the specs refer to
  - you are expect to read the **grading guidelines**
  - it's your responsibility



# Programming & Good Habbits

➡ ***Always*** check return code!


- ▬ `open()`, `write()`
- ▬ `malloc()`
- ▬ `switch (errno) { ... }`

➡ Initialize ***all*** variables!

- ▬ `int i=0;`
- ▬ `char *p=NULL;`
- ▬ `struct timeval timeout;`  
`memset(&timeout, 0, sizeof(struct timeval));`

➡ ***Never*** leak any resources!

- ▬ `malloc()` and `free()`
- ▬ `open()` and `close()`
- ▬ delete temporary files



# Programming & Good Habbits

➡ ***Don't*** assume external input will be short

- use `strncpy()` and not `strcpy()`
- use `snprintf()` and not `sprintf()`
- use `sizeof()` and not a constant, for example,

```
unsigned char buf[80];
```


```
buf[0] = '\0'; /* initialization */
```

```
strncpy(buf, argv[1], sizeof(buf));
```

```
buf[sizeof(buf)-1] = '\0'; /* in case argv[1] is long */
```

➡ Fix your code so that you have ***zero*** compiler warnings!

- use `-Wall` when you compile to get all compiler warnings



# Notes on gdb

➡ The debugger is your friend! Get to know it **NOW!**

compile program with: `-g`

start debugging: `gdb [-tui] listtest`

set breakpoint: `(gdb) break main`

`(gdb) break listtest.c:87`

run program (w/ arguments): `(gdb) run [arg1 arg2 ...]`

clear breakpoint: `(gdb) clear`

stack trace: `(gdb) where`

print field: `(gdb) print pList->anchor`

print in hex: `(gdb) print/x pList->anchor`

single-step at same level: `(gdb) next`


single-step into a function: `(gdb) step`

print field after every cmd: `(gdb) display pList->num_members`

assignment: `(gdb) set pList->num_members=99`

continue: `(gdb) cont`

quit: `(gdb) quit`




# Warmup #1




2 parts

- develop a *doubly-linked circular list* called *My402List*
  - this corresponds to part (A) of the grading guidelines
  - to implement a traditional *linked-list abstraction*
    - ◆ internally, the implementation is a *circular list*
    - ◆ internally, it behaves like a *traditional list*
    - ◆ why? circular list implementation may be a little "cleaner"
- use your doubly-linked circular list to implement a command:
  - *sort* - sort a list of bank transactions
  - this corresponds to part (B) of the grading guidelines




# A Linked-List Abstraction

- ➡ A list of elements, linked so that you can move from one to the next (and/or previous)
  - ▬ each element holds an object of some sort
- ➡ *Functionally:*
  - ▬ First()
  - ▬ Next()
  - ▬ Last()
  - ▬ Prev()
  - ▬ Insert()
  - ▬ Remove()
  - ▬ Count()
- ➡ Need to have a well-defined interface
  - ▬ once you have a good interface, if the implementation is broken, fix the implementation!
    - don't fix the "application"


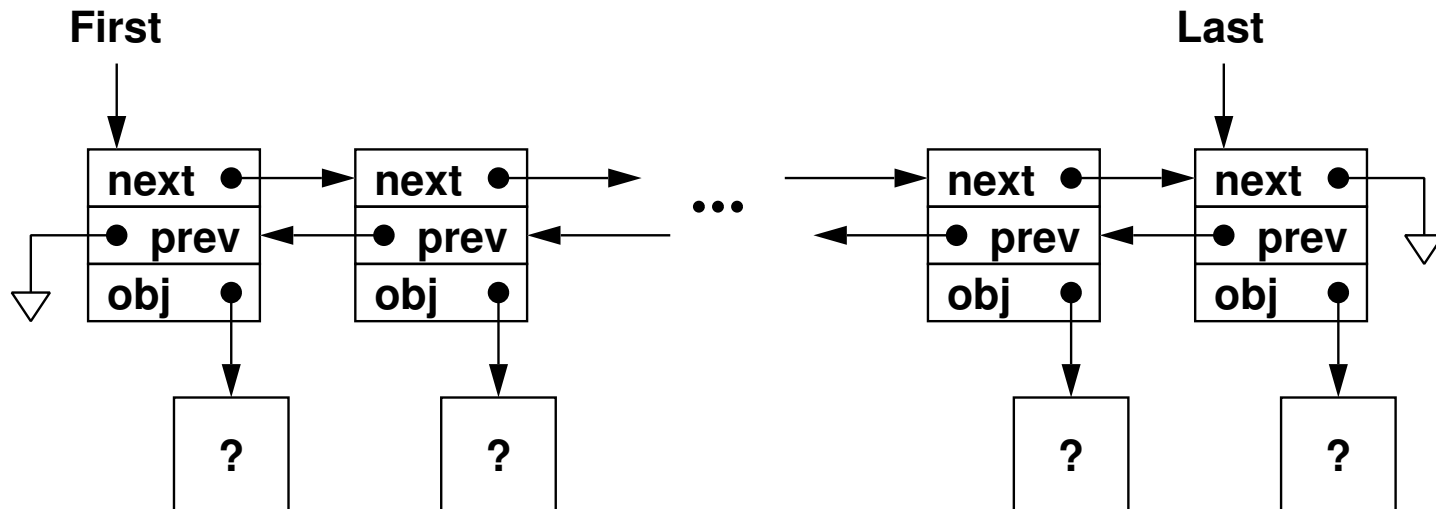


# A Linked-List Abstraction

- ➡ There are basically two types of lists
  - 1) next/prev pointers in object
  - 2) next/prev pointers outside of object
- ➡ (1) has a major drawback that a list item cannot be inserted into multiple lists



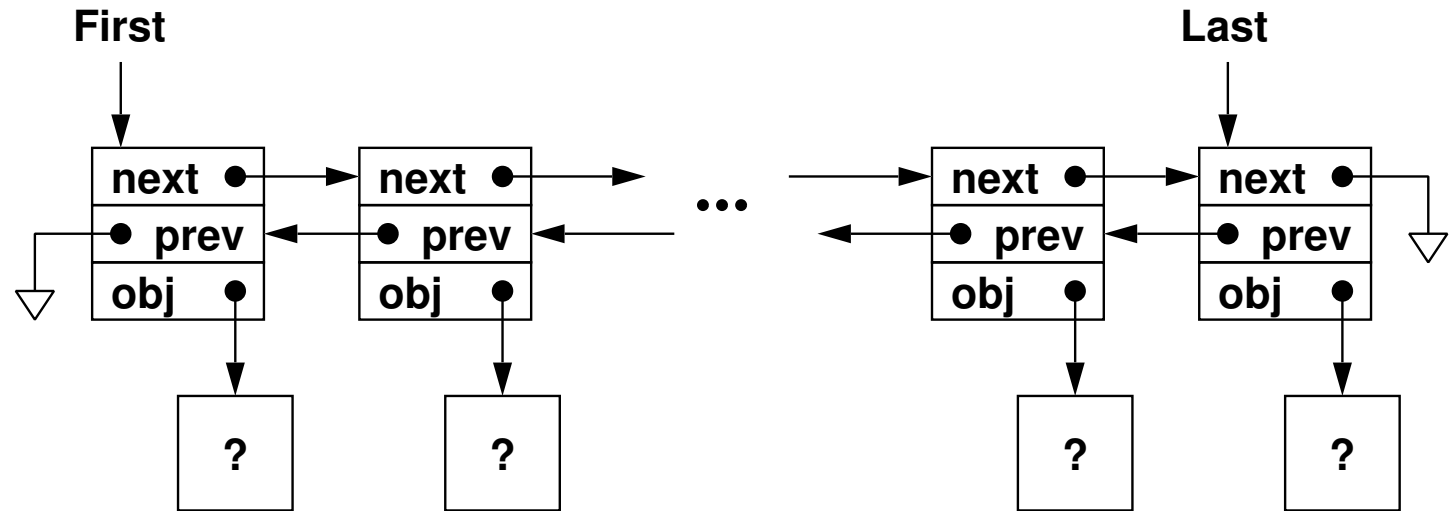
— we will implement (2) in warmup1, our kernel uses (1)





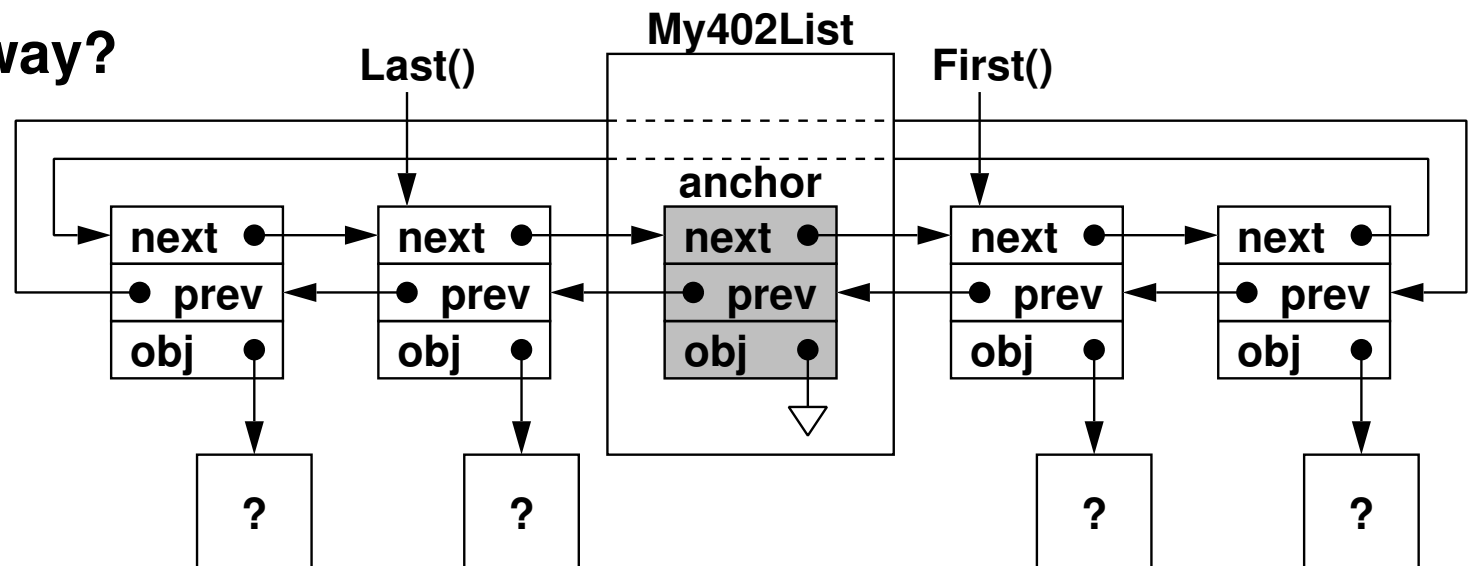
# Doubly-linked Circular List

## ➡ Abstraction




## ➡ Implementation

— why this way?



— your job is to implement the traditional list abstraction  
using a circular list



# my402list.h

```
#ifndef _MY402LIST_H_
#define _MY402LIST_H_

#include "cs402.h"


typedef struct tagMy402ListElem {
    void *obj;
    struct tagMy402ListElem *next;
    struct tagMy402ListElem *prev;
} My402ListElem;

typedef struct tagMy402List {
    int num_members;
    My402ListElem anchor;

    /* You do not have to set these function pointers */
    int (*Length)(struct tagMy402List *);
    int (*Empty)(struct tagMy402List *);

    int (*Append)(struct tagMy402List *, void*);
    int (*Prepend)(struct tagMy402List *, void*);
    void (*Unlink)(struct tagMy402List *, My402ListElem*);
    void (*UnlinkAll)(struct tagMy402List *);
};
```

➡ You need to learn to ignore things you don't understand  
— assume that they are perfect



# my402list.h

```

int  (*InsertBefore)(struct tagMy402List *, void*, My402ListElem*);
int  (*InsertAfter)(struct tagMy402List *, void*, My402ListElem*);

My402ListElem *(*First)(struct tagMy402List *);
My402ListElem *(*Last)(struct tagMy402List *);
My402ListElem *(*Next)(struct tagMy402List *, My402ListElem *cur);
My402ListElem *(*Prev)(struct tagMy402List *, My402ListElem *cur);

My402ListElem *(*Find)(struct tagMy402List *, void *obj);
} My402List;

extern int  My402ListLength(My402List*);
extern int  My402ListEmpty(My402List*);

extern int  My402ListAppend(My402List*, void*);
extern int  My402ListPrepend(My402List*, void*);
extern void My402ListUnlink(My402List*, My402ListElem*);
extern void My402ListUnlinkAll(My402List*);
extern int  My402ListInsertAfter(My402List*, void*, My402ListElem*);
extern int  My402ListInsertBefore(My402List*, void*, My402ListElem*);

extern My402ListElem *My402ListFirst(My402List*);
extern My402ListElem *My402ListLast(My402List*);
extern My402ListElem *My402ListNext(My402List*, My402ListElem*);
extern My402ListElem *My402ListPrev(My402List*, My402ListElem*);


extern My402ListElem *My402ListFind(My402List*, void*);

extern int  My402ListInit(My402List*);
#endif /*_MY402LIST_H_*/

```



**You need to implement all the mentioned functions**



# my402list.c

## ➡ How to start?

- ▬ `cp my402list.h my402list.c`
- ▬ edit `my402list.c` in a text editor
  - replace data structure declarations with `"#include"`
  - change all function declarations to function implementations
    - ◆ remove `"extern"` and implement function

```
#include "my402list.h"


extern int  My402ListLength(My402List*);
extern int  My402ListEmpty(My402List*);

extern int  My402ListAppend(My402List*, void*);
extern int  My402ListPrepend(My402List*, void*);
extern void My402ListUnlink(My402List*, My402ListElem*);
extern void My402ListUnlinkAll(My402List*);
extern int  My402ListInsertAfter(My402List*, void*, My402ListElem*);
extern int  My402ListInsertBefore(My402List*, void*, My402ListElem*);

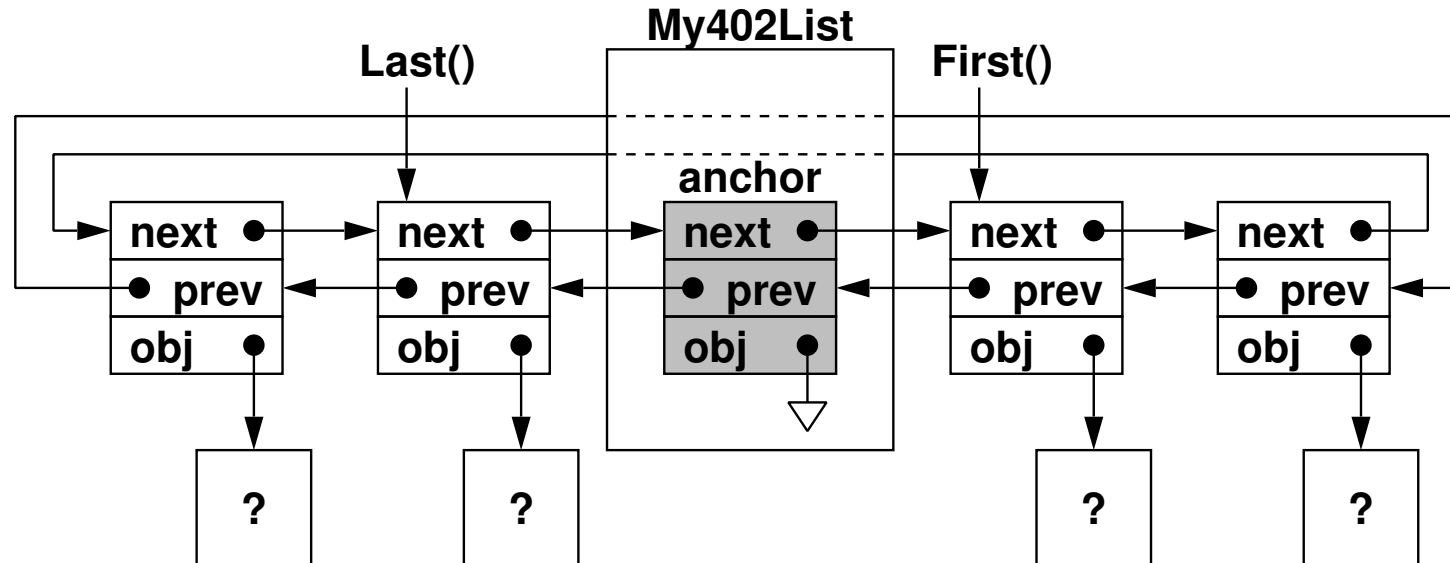
extern My402ListElem *My402ListFirst(My402List*);
extern My402ListElem *My402ListLast(My402List*);
extern My402ListElem *My402ListNext(My402List*, My402ListElem*);
extern My402ListElem *My402ListPrev(My402List*, My402ListElem*);

extern My402ListElem *My402ListFind(My402List*, void*);

extern int  My402ListInit(My402List*);
```



# Implementation




```
int Length() { return num_members; }
int Empty() { return num_members<=0; }
```

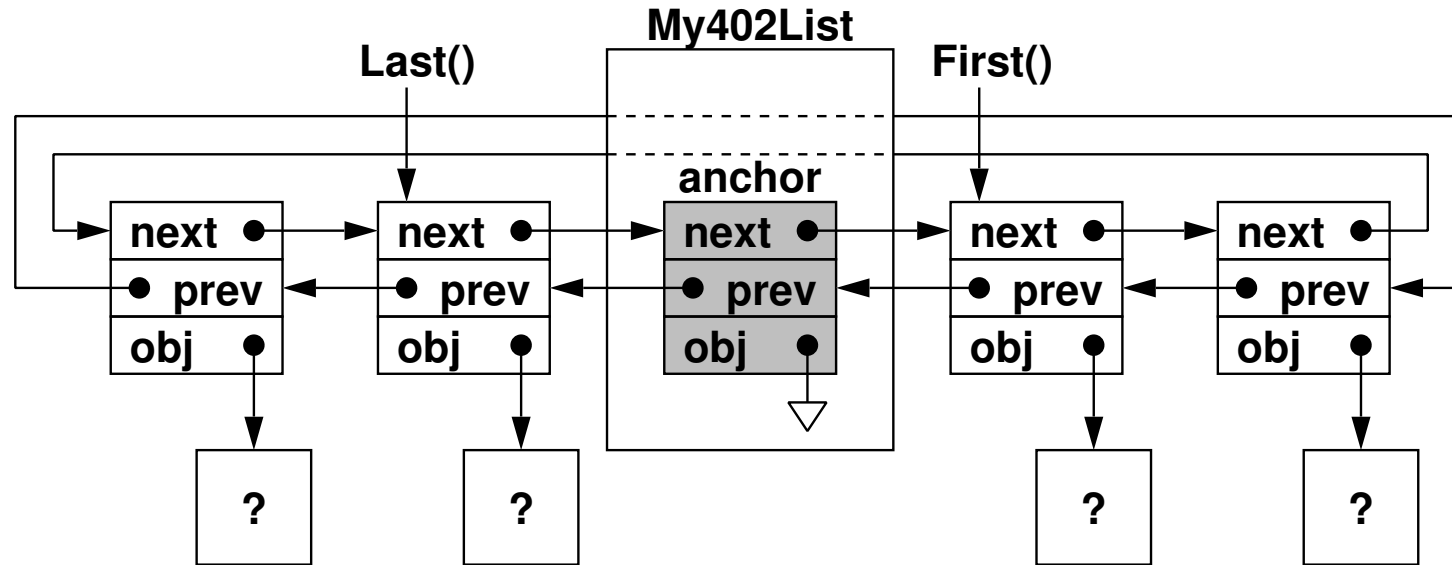
```
int Append(void *obj);
int Prepend(void *obj);
void Unlink(My402ListElem*);
void UnlinkAll();
int InsertBefore(void *obj, My402ListElem *elem);
int InsertAfter(void *obj, My402ListElem *elem);
```

```
My402ListElem *First();
My402ListElem *Last();
My402ListElem *Next(My402ListElem *cur);
My402ListElem *Prev(My402ListElem *cur);
```

```
My402ListElem *Find(void *obj);
int Init();
```



# Usage - Traversing the List




```
void Traverse(My402List *list)
{
    My402ListElem *elem=NULL;

    for (elem=My402ListFirst(list);
        elem != NULL;
        elem=My402ListNext(list, elem)) {
        Foo *foo=(Foo*)(elem->obj);

        /* access foo here */
    }
}
```


➡ This is how an *application* will use **My402List**

- ➡ you must support this "*contract*" with you application
- ➡ if broken, fix the "implementation" and not the "application"




# listtest

- ➡ Use provided `listtest.c` and `Makefile` to create `listtest`
  - ▬ `listtest` must run without error and you must not change `cs402.h`, `my402list.h`, `listtest.c` and `Makefile`
  - ▬ they specify how your code in `my402list.c` is expected to be used
- ➡ You should learn how to run `listtest` under `gdb`



# Sort Command

- ➡ **warmup1 sort [tfile]**
  - produce a sorted transaction history for the transaction records in `tfile` (or `stdin`) and compute balances
- ➡ **Input is an ASCII text file (use `fgets()` to read a line)**
  - each line in a `tfile` contains 4 fields delimited by `<TAB>`
    - transaction type (single character)
      - ◆ "+" for deposit
      - ◆ "-" for withdrawal
    - transaction time (UNIX time)
      - ◆ `man -s 2 time`
    - amount (a number, a period, two digits)
    - transaction description (textual description)
      - ◆ cannot be empty
- ➡ **Output must be in the specified format exactly**
  - use the grading guidelines to check if you miss something
    - formatting bugs should be very easy to fix





# Sort Command

## ➡ Output


```
00000000011111111122222222233333333334444444445555555556666666667777777778
1234567890123456789012345678901234567890123456789012345678901234567890
```

Date	Description	Amount	Balance
Thu Aug 21 2008	...	1,723.00	1,723.00
Wed Dec 31 2008	...	( 45.33)	1,677.67
Mon Jul 13 2009	...	10,388.07	12,065.74
Sun Jan 10 2010	...	( 654.32)	11,411.42

## ➡ How to keep track of balance

- first thing that comes to mind is to use `double`
- the weird thing is that if you are not very careful with `double`, your output will be wrong (by 1 penny) once in a while
- recommendation: keep the balance in cents, not dollars
  - no precision problem with integers!

## ➡ Read *grading guidelines* and find many examples of *valid input* and *expected printout*



# Sort Command


```
00000000011111111112222222222333333333344444444445555555555666666666677777777778
12345678901234567890123456789012345678901234567890123456789012345678901234567890
```

Date	Description	Amount	Balance
Thu Aug 21 2008	...	1,723.00	1,723.00
Wed Dec 31 2008	...	( 45.33)	1,677.67
Mon Jul 13 2009	...	10,388.07	12,065.74
Sun Jan 10 2010	...	( 654.32)	11,411.42

➡ The spec requires you to call `ctime()` to convert a Unix timestamp to string

- ➡ then pick the right characters to display as date
- ➡ e.g., `ctime()` returns "Thu Aug 30 08:17:32 2012\n"
- be careful, `ctime()` returns a pointer that points to a *global variable*, so you must *make a copy*

```
char date[16];
char buf[26];
strncpy(buf, ctime(...), sizeof(buf));
date[0] = buf[0];
date[1] = buf[1];
...
date[15] = '\0';
```



# Sort Command


```
00000000011111111112222222222333333333344444444445555555555666666666677777777778
12345678901234567890123456789012345678901234567890123456789012345678901234567890
```

Date	Description	Amount	Balance
Thu Aug 21 2008	...	1,723.00	1,723.00
Wed Dec 31 2008	...	( 45.33)	1,677.67
Mon Jul 13 2009	...	10,388.07	12,065.74
Sun Jan 10 2010	...	( 654.32)	11,411.42




## Format your data in your own buffer

- write a function to "format" numeric fields into null-terminated strings
  - it's a little more work, but you really should have this code isolated
    - ◆ in case you have bugs, just fix this function
- you can even do the formatting when you append or insert your data structure to your list
  - need more fields in your data structure
- this way, you can just print things out easily
- use `printf("%s", ...)` to print a field to `stdout`




# Warmup #1

- ➡ I'm giving you a lot of details on how to do things in C
  - this is the first and last assignment that I will do this!
  - you must learn C (and Unix) on your own
- ➡ Read man pages
- ➡ Ask questions in class Google Group
  - or send e-mail to me
- ➡ Come to office hours, especially if you are stuck




# Some General Requirements

- ➡ Some major requirements for all programming assignments
- severe penalty for failing `make` (can lose up to 10 points)
    - we will attempt to fix your Makefile if we cannot compile your code
    - we are *not permitted to change your code*
  - severe penalty for using large memory buffers
    - if input file is large, you must not read the whole file into into a large memory buffer
      - ◆ must learn how to read a large file properly
  - severe penalty for any segmentation fault -- you must test your code well
  - severe penalty for not using separate compilation or for having all your source code in header files -- you must learn to plan how to write your program
    - read warmup1 FAQ to see what's the best way to go about this




# Grading Requirements

- ➡ For *warmup assignments*, it's important that **every byte** of your data is read and written correctly
  - you are not entitled to partial credit just because you wrote some code
- ➡ For *warmup assignments*, you should run your code against the *grading guidelines* on 32-bit Ubuntu 16.04
  - must not change the commands there
    - we will change the data for actual grading, but we will stick to the commands (as much as we can)
  - to be fair to all, running scripts in the grading guidelines on the grader's 32-bit Ubuntu 16.04 is *the only way we can grade*




# Separate Compilation

- ➡ Break up your code into *modules*
  - ➡ *compile the modules separately*, at least one rule per module per rule in the `Makefile`
  - ➡ a separate rule to *link* all the modules together
    - if your program requires additional libraries, add them to the link stage
- ➡ To receive full credit for separate compilation
  - ➡ to create an executable, at a minimum, you must run the compiler at least *twice* and the linker *once*
  - ➡ see the *warmup1 FAQ* for exactly how to avoid losing points



# README


- ➡ Start with the README templates from the spec
  - **BUILD & RUN** (required)
    - replace "(Comments: ?)" with appropriate responses
  - **SELF-GRADING** (required)
    - replace each "?" with a **numerical score**
  - **BUGS / TESTS TO SKIP** (required)
    - replace "(Comments: ?)" with a list of tests to skip or "none"
      - ◆ you would still lose points, but this may prevent losing additional points in another part
  - **ADDITIONAL INFORMATION FOR GRADER** (optional)
    - grader must read this
  - **OTHERS** (optional)
    - will **not** be considered for grading
- ➡ There should be no "?" left in a response **in a required section** after you have filled out a README file correctly
  - 0.5 pt will be deducted if a "?" is not replaced with something appropriate or if the line is omitted






# Code Design - Functional vs. Procedural

- ➡ Don't design your program "procedurally"
- ➡ You need to learn how to write functions!
  - please note that this is not "functional programming" ("functional programming" is something else)
  - a function has a well-defined interface
    - what are the meaning of the parameters
    - what does it suppose to return
  - pre-conditions
    - what must be true when the function is entered
    - you assume that these are true
      - ◆ you can verify it if you want
  - post-conditions
    - what must be true when the function returns
  - you design your program by making designing a sequence of function calls




# Warmup #1 - Miscellaneous Requirements

- ➡ Run your code against the *grading guidelines*
  - must not change the test program
- ➡ You must not use any *external code fragments*
- ➡ You must not use *array* to implement any list functions
  - *must use pointers* because this is a pointer exercise
  - read my *review about pointers* in warmup1 FAQ
- ➡ It's important that every byte of your data is read and written correctly.
  - `diff` commands in the grading guidelines must *not* produce *any* output or you will not get credit
    - what does "not produce any output" mean?
      - ◆ it means exactly what it says!
- ➡ Please see Warmup #1 spec for additional details
  - please read the *entire* spec (including the grading guidelines) *yourself*



# Development


- ➡ Get familiar with "*Warmup #1 FAQ*" and "*Programming FAQ*"
- ➡ Text Editors
  - emacs, pico, vi
  - some students like Sublime Text
    - you are on your own with Sublime Text
- ➡ Compiler
  - "gcc --version" should say it's version 5.4.something
- ➡ IDE
  - some students like Eclipse
    - you are on your own with Eclipse
  - the grader is *not permitted* to use an IDE to compile or run your program
    - if you use an IDE, it's your responsibility to make sure that you provide a Makefile so that the grader can type the command in the spec to compile



# Warmup #1 (Part 2)


Bill Cheng

*<http://merlot.usc.edu/william/usc/>*



# listtest

- ➡ Use provided `listtest.c` and `Makefile` to create `listtest`
  - ▬ `listtest` must run without error and you must not change `listtest.c` and `Makefile`
  - ▬ They specifies how your code is expected to be used
- ➡ You should learn how to run `listtest` under `gdb`



# gdb listtest Exercise



Do the following gdb exercise with `listtest`

- = **IMPORTANT:** draw picture on a piece of paper!
- = first, change "`num_items=64`" in `DoTest()` to "`num_items=3`"

make

`gdb listtest`

`(gdb) break DoTest`

`(gdb) run`

`(gdb) n`      ← do this 6 times, you are now at call to `CreateTestList()`

`(gdb) print &(list.anchor)`      ← what's the address of the anchor?

`(gdb) print list`      ← does the list look like an empty list?

`(gdb) n`      ← returned from `CreateTestList()`

`(gdb) print list`      ← does the list look like a 3-item list?


`(gdb) print list.anchor`      ← what's in the anchor?

`(gdb) print *(list.anchor.next)`

`(gdb) print *(list.anchor.next->next)`

`(gdb) print *(list.anchor.next->next->next)`

this should be the last list element,  
does its `next` pointer point to the anchor?




# C Review: C-string Functions

➡ C-string manipulating functions in your kernel assignments (some do not assume C-string):

```
int      memcmp(const void *cs, const void *ct, size_t count);
void     *memcpy(void *dest, const void *src, size_t count);
int      strncmp(const char *cs, const char *ct, size_t count);
int      strcmp(const char *cs, const char *ct);
char     *strcpy(char *dest, const char *src);
char     *strncpy(char *dest, const char *src, size_t count);
void     *memset(void *s, int c, size_t count);
size_t   strlen(const char *s);
size_t   strnlen(const char *s, size_t count);
char     *strchr(const char *s, int c);
char     *strrchr(const char *s, int c);
char     *strstr(const char *s1, const char *s2);
char     *strcat(char *dest, const char *src);
char     *strdup(const char *s);
char     *strtok(char *s, const char *d);
```

⇒ you also need:


```
char     *fgets(char *s, int size, FILE *stream);
int      printf(const char *format, ...);
```



# C File I/O Review: Reading Text Input

- ➡ Read in an entire line using `fgets()`
  - especially since we know the maximum line length
- ➡ If a filename is given, use `fopen()` to get a file pointer (`FILE*`)
  - `FILE *fp = fopen(..., "r");`
  - read man pages of `fopen()`
  - if a filename is not given, you will be reading from "standard input" (i.e., file descriptor 0)
  - `FILE *fp = stdin;`
  - see grading guidelines for examples
  - `cat ... | ./warmup1 sort`
  - pass the file pointer around so that you run the same code whether you input comes from a file or `stdin`

```
My402List list;
if (!My402ListInit(&list)) { /* error */ }
if (!ReadInput(fp, &list)) { /* error */ }
if (fp != stdin) fclose(fp);
SortInput(&list);
PrintStatement(&list);
```





# C File I/O Review: Parsing Text Input


## ➡ Read a line

```
char buf[2000];
if (fgets(buf, sizeof(buf), fp) == NULL) {
    /* end of file */
} else {
    /* parse it */
}
```

## ➡ Parse a line according to the spec

- find an *easy* and *correct* way to parse the line
  - according to the spec, each line must have exactly 3 <TAB> characters
  - I think it's easy and correct to go after this

```
char *start_ptr = buf;
char *tab_ptr = strchr(start_ptr, '\t');
if (tab_ptr != NULL) {
    *tab_ptr++ = '\0';
}
/* start_ptr now contains a
   "null-terminated string" */
```



# C File I/O Review: Parsing Text Input

```

→ char *start_ptr = buf;
   char *tab_ptr = strchr(start_ptr, '\t');
   if (tab_ptr != NULL) {
       *tab_ptr++ = '\0';
   }
   /* start_ptr now contains a
      "null-terminated string" */


```

buf

'a'	'b'	'c'	'\t'	' '	'd'	'e'	'\t'	'f'	'\0'	'\0'
-----	-----	-----	------	-----	-----	-----	------	-----	------	------



start\_ptr

tab\_ptr




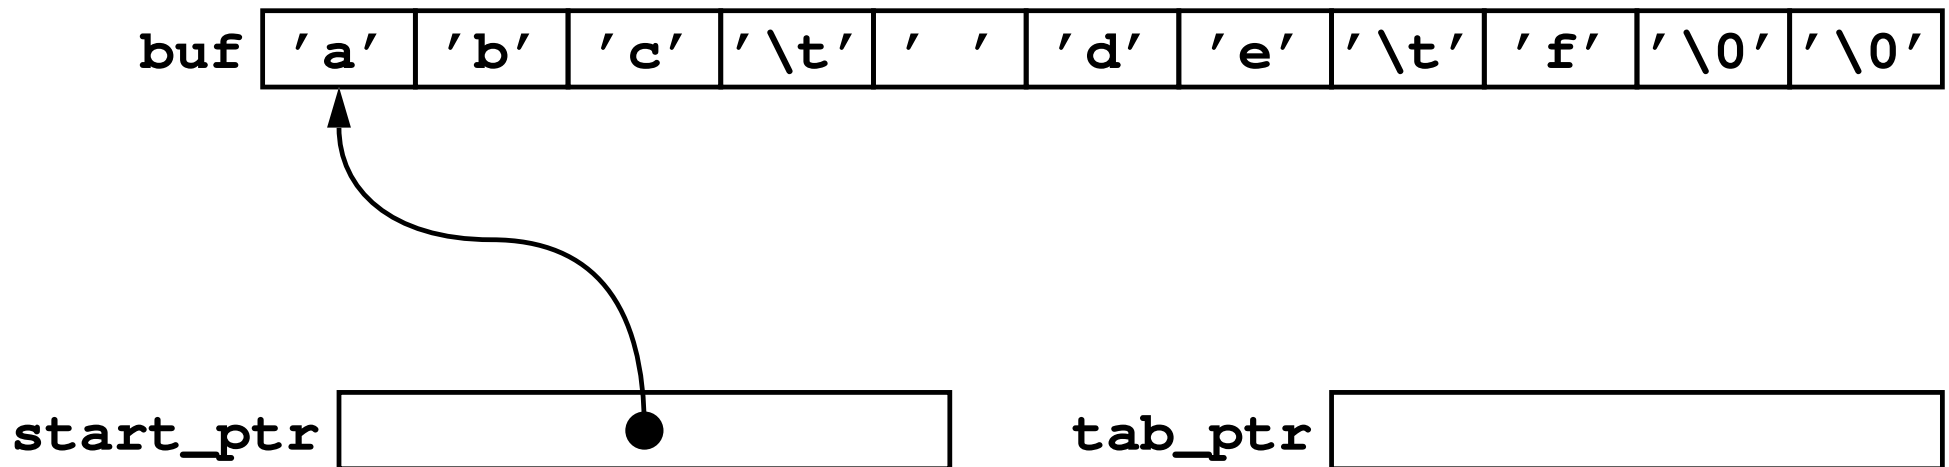
# C File I/O Review: Parsing Text Input

```
➔ char *start_ptr = buf;  
char *tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```




# C File I/O Review: Parsing Text Input

```
char *start_ptr = buf;  
➔ char *tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```



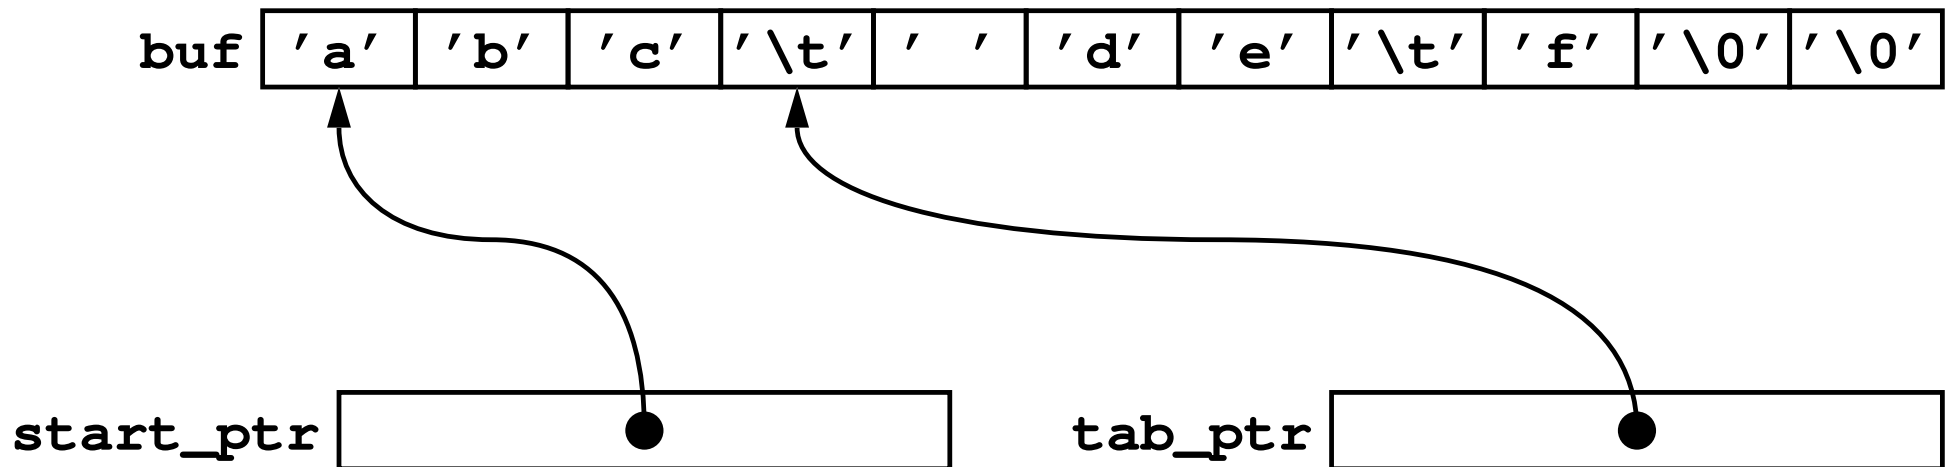
# C File I/O Review: Parsing Text Input

```
char *start_ptr = buf;  
➔ char *tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```



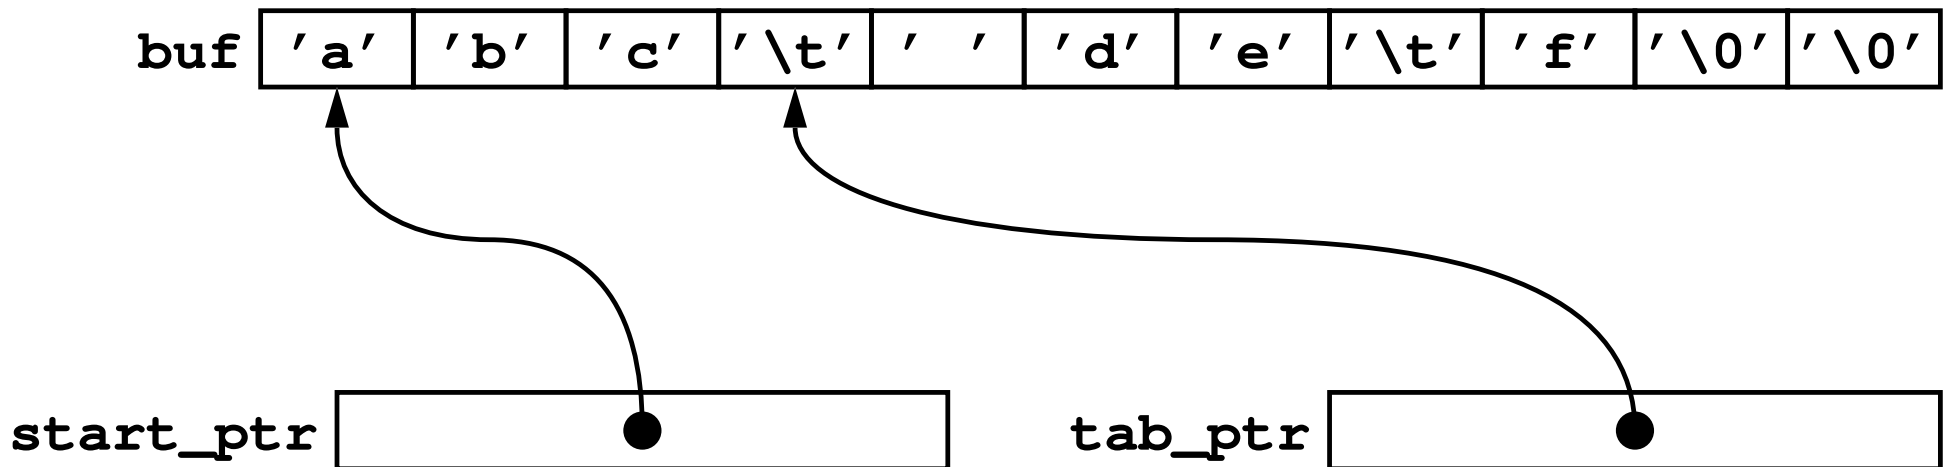
# C File I/O Review: Parsing Text Input

```
char *start_ptr = buf;  
char *tab_ptr = strchr(start_ptr, '\t');  
→ if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```



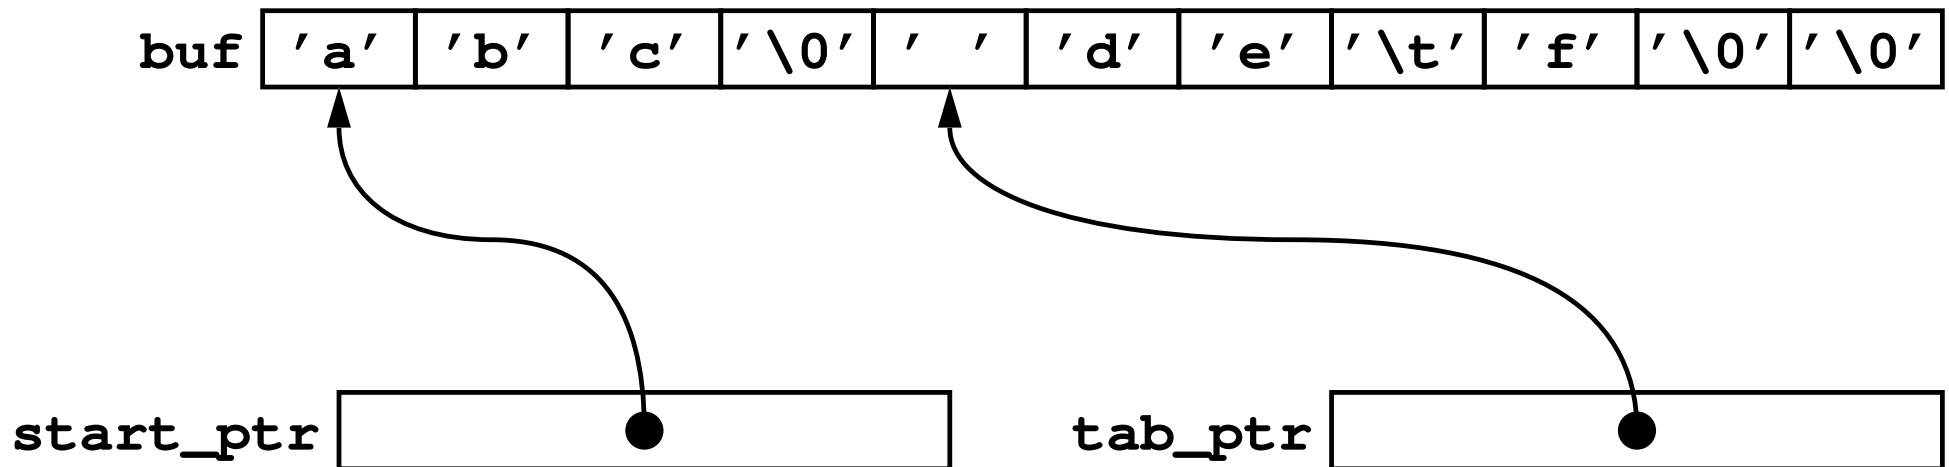
# C File I/O Review: Parsing Text Input

```
char *start_ptr = buf;  
char *tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
→   *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```



# C File I/O Review: Parsing Text Input


```
char *start_ptr = buf;  
char *tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
→   *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```






# C File I/O Review: Parsing Text Input (2nd Iteration)

```
➔ start_ptr = tab_ptr;  
tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```




# C File I/O Review: Parsing Text Input (2nd Iteration)

```
start_ptr = tab_ptr;  
➔ tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```




# C File I/O Review: Parsing Text Input (2nd Iteration)

```
start_ptr = tab_ptr;  
tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
→ *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```




# C File I/O Review: Parsing Text Input (2nd Iteration)

```
start_ptr = tab_ptr;  
tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
→ *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```




# C File I/O Review: Parsing Text Input (3rd Iteration)

```
➔ start_ptr = tab_ptr;  
tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```



# C File I/O Review: Parsing Text Input (3rd Iteration)

```
start_ptr = tab_ptr;  
➔ tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```



# C File I/O Review: Parsing Text Input (3rd Iteration)

```
start_ptr = tab_ptr;  
➔ tab_ptr = strchr(start_ptr, '\t');  
if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```


buf

'a'	'b'	'c'	'\0'	' '	'd'	'e'	'\0'	'f'	'\0'	'\0'
-----	-----	-----	------	-----	-----	-----	------	-----	------	------

start\_ptr



tab\_ptr



# C File I/O Review: Parsing Text Input (3rd Iteration)

```
start_ptr = tab_ptr;  
tab_ptr = strchr(start_ptr, '\t');  
→ if (tab_ptr != NULL) {  
    *tab_ptr++ = '\0';  
}  
/* start_ptr now contains a  
   "null-terminated string" */
```


buf

'a'	'b'	'c'	'\0'	' '	'd'	'e'	'\0'	'f'	'\0'	'\0'
-----	-----	-----	------	-----	-----	-----	------	-----	------	------

start\_ptr



tab\_ptr






# C File I/O Review: Parsing Text Input (3rd Iteration)

```


start_ptr = tab_ptr;
tab_ptr = strchr(start_ptr, '\t');
if (tab_ptr != NULL) {
    *tab_ptr++ = '\0';
}
/* start_ptr now contains a
   "null-terminated string" */

```

buf [ 'a' | 'b' | 'c' | '\0' | ' ' | 'd' | 'e' | '\0' | 'f' | '\0' | '\0' ]



- found only 2 tabs
- need to keep a count



# C Function Review: Returning A C-String

- ➡ How to return a C-string?
- easiest way is to provide a buffer in the caller
    - can only do this if you know what buffer size to use


- ➡ Caller code (to get a C-string from callee):

```
char field[1024];

if (GetField(field, sizeof(field) > 0) {
    /* field now has a C-string */
} else {
    /* print error message */
    exit(-1);
}
```


- ➡ Callee code:

```
int GetField(char *buf, int buf_sz)
{
    /* setup start_ptr to point to a field */
    strncpy(buf, start_ptr, buf_sz);
    buf[buf_sz-1] = '\0';
    return strlen(buf);
}
```




# Warmup #1

- ➡ I'm giving you a lot of details on how to do things in C
  - this is the first and last assignment that I will do this!
  - you must learn C on your own
- ➡ Read man pages
- ➡ Ask questions in class Google Group
  - or send e-mail to me
- ➡ Come to office hours, especially if you are stuck



# Warmup #1 - Miscellaneous Requirements

- ➡ Run your code against the *grading guidelines*
  - must not change the test program
- ➡ You must not use any *external code fragments*
- ➡ You must not use *array* to implement any list functions
  - must use pointers
- ➡ If input file is large, you must not read the whole file into into a large memory buffer
- ➡ It's important that every byte of your data is read and written correctly.
  - `diff` commands in the grading guidelines must **not** produce **any** output or you will not get credit
- ➡ Please see Warmup #1 spec for additional details
  - please read the *entire* spec *yourself*



# Warmup #1 - Miscellaneous Requirements

➡ Explain how Makefile work

```
listtest: listtest.o my402list.o
        gcc -o listtest -g listtest.o my402list.o


listtest.o: listtest.c my402list.h
        gcc -g -c -Wall listtest.c

my402list.o: my402list.c my402list.h
        gcc -g -c -Wall my402list.c

clean:
        rm -f *.o listtest
```

➡ Explain why the above Makefile satisfied *separate compilation* requirement

➡ Explain how to convert the above Makefile to a Makefile students can use for part (2)



# Demos

- ➡ If there is time in discussion section, demonstrate how create a `warmup1` subdirectory on 32-bit Ubuntu 16.04
  - if there is no time in discussion section, students should try all these on their own


- ➡ Create "hello.c" on 32-bit Ubuntu 16.04

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    printf("Hello World!\n");
    return 0;
}
```

- ➡ Demonstrate how to run gcc

```
gcc -g -Wall hello.c
```



# Demos

➡ Create "args.c" on 32-bit Ubuntu 16.04


```
#include <stdio.h>

static int foobar = 0x12345678;

int main(int argc, char *argv[])
{
    for (int i = 0; i < argc; i++) {
        printf("argv[%d] = '%s'\n", i, argv[i]);
    }
    return 0;
}
```

➡ Compile and run:

```
gcc -g -Wall args.c
./a.out x y z
```




# Demos

➡ Demonstrate how to debug a.out

```
gdb a.out
(gdb) break main
(gdb) run abc xyz
(gdb) print argc
(gdb) print argv[0]
(gdb) print argv[1]
(gdb) print argv[2]
(gdb) print &argc
(gdb) print &i
(gdb) print &foobar
(gdb) next
(gdb)
```

➡ Go to warmup1 spec and demonstrate how to copy and paste from the grading guidelines

➡ Go to warmup1 spec and demonstrate how to make a submission and get a PIN





# Unix Commands

➡ Walk through and demonstrate the commands on the Unix Command Line Reference web page

- click on the "summary of some commonly used Unix commands" link at the bottom of the class home page

`ls`

`cat`

`more`

`pwd`

`mkdir (directory name)`

`cd`

`cp (src file path) (dest file path)`

`mv (src file path) (dest file path)`

`man (cmd name)`

`rm (file path)`


`rmdir (empty directory name)`

`ps`

`kill (proc id)`

`pico (file path)`


`exit`



# Warmup #2 (Part 1)

Bill Cheng

*<http://merlot.usc.edu/william/usc/>*



# Multi-threading Exercise



Make sure you are familiar with the *pthread*s library

= Ch 2 of textbook - threads, signals

- additional resource is a book by Nichols, Buttlar, and Farrell  
“*Pthreads Programming*”, O’Rielly & Associates, 1996

= you must learn how to use pthreads *mutex* and *condition variables* correctly

- `pthread_mutex_lock()` / `pthread_mutex_unlock()`

- `pthread_cond_wait()` / `pthread_cond_broadcast()`

- ◆ do *not* use `pthread_cond_signal()` for warmup2

= you must learn how to handle UNIX *signals* (<Ctrl+C>)


- `sigprocmask()` / `sigwait()`

- `pthread_cancel()`


= you may want to learn how to disable/enable *cancellation* in pthreads

- `pthread_setcancelstate()`

next week



# Token Bucket Filter




**Ex:**

— traffic controller/shaper




Your job is to implement 4 *cooperating child threads* to move the packets along by following rules described in the spec

— the main thread creates these threads, join with them, then print statistics





# We Are *Not* Doing Event-driven Simulation

- ➡ An *event queue* is a sorted list of events according to timestamps; smallest timestamp at the head of queue
  - event has zero duration (events can happen at the same time)
- ➡ *Object oriented*: every object has a "next event" (what and when it will do next), this event is inserted into the event queue
- ➡ Execution: remove an event from the head of queue, "execute" the event (notify the corresponding object so it can insert the next event)
- ➡ Insert into the event queue according to timestamp of a new event; insertion may cause additional events to be deleted or inserted
- ➡ Potentially repeatable runs (if the same seed is used to initialize random number generator)
- ➡ The simulator never "sleeps"; it tries to run as fast as it can to finish the simulation as quickly as possible




# We Are *Not* Doing Event-driven Simulation

- ➡ Multiple event can happen at the same time in an *event-driven simulation*
- we will *not* be doing that!




$$r_3 = d_3 - a_3$$




# "Time Driven" Simulation


- ➡ We will use the words "simulation" and "emulation" interchangeably
- ➡ No "event queue"
  - every *active object* is implemented as a *thread*
  - threads interacting with one another through the use of *shared variables*
    - how else can threads "talk" to each other?!
- ➡ It takes time to execute simulation code
  - the time it takes to do all that is part of the simulation
  - to simulation the passing of time, call `usleep()`
    - e.g., if doing something takes  $x$  usec, call `usleep(x)`
    - Ubuntu does not run a "realtime" OS, it's "best effort"
    - `usleep(x)` will return more than  $x$  usec later
      - ◆ and sometimes, *a lot more* than  $x$  usec later
      - ◆ you need to decide if the extra delay is reasonable or it's due to a bug in your code



# "Time Driven" Simulation





- ➡ Let your machine decide which thread to run next
  - results can never be reproducible exactly
  - debugging can be more challenging
- ➡ Compete for resources (such as  $Q1$ ,  $Q2$ , and anything shared), must use a *single mutex*
- ➡ *No busy-waiting*
  - must use a *single CV*






# Arrivals & Departures

- ▢  $a_i$  : arrival time
- ▢  $d_i$  : departure time
- ▢  $s_i$  : service time
- ▢  $r_i$  : response time (time in system)
- ▢  $q_i^1, q_i^2$  : queueing/waiting time





▢  $r_1 = d_1 - a_1$




# Arrivals & Departures

- ▢  $a_i$  : arrival time
- ▢  $d_i$  : departure time
- ▢  $s_i$  : service time
- ▢  $r_i$  : response time (time in system)
- ▢  $q_i^1, q_i^2$  : queueing/waiting time





▢  $r_2 = d_2 - a_2$




# Arrivals & Departures

- ▢  $a_i$  : arrival time
- ▢  $d_i$  : departure time
- ▢  $s_i$  : service time
- ▢  $r_i$  : response time (time in system)
- ▢  $q_i^1, q_i^2$  : queueing/waiting time



▢  $r_3 = d_3 - a_3$



# Simulation/Emulation



## Two simulation modes



- 1) **Deterministic**: fixed inter-arrival time ( $1/\lambda$ ), token requirement ( $P$ ), and service time ( $1/\mu$ )
- 2) **Trace-driven**: every packet has its own *inter-arrival time*, *token requirement*, and *service time* (a line in a "tsfile")

— if you think about it


carefully, there is really no

difference between these two modes

- write your code for the trace-driven mode
- if running in deterministic mode, instead of reading a line from the "tsfile" to create a packet, just create a packet using information stored in global variables




# Simulation/Emulation



➡ You will need to implement 4 *cooperating child threads*

- ➡ *packet* arrival thread
- ➡ *token* depositing thread
- ➡ two *server* threads
- ➡ these threads work together to simulate the operation of this token bucket filter
  - threads work together using *shared variables*



# Simulation/Emulation

➡ Very high level pseudo-code for the *packet/token thread*:


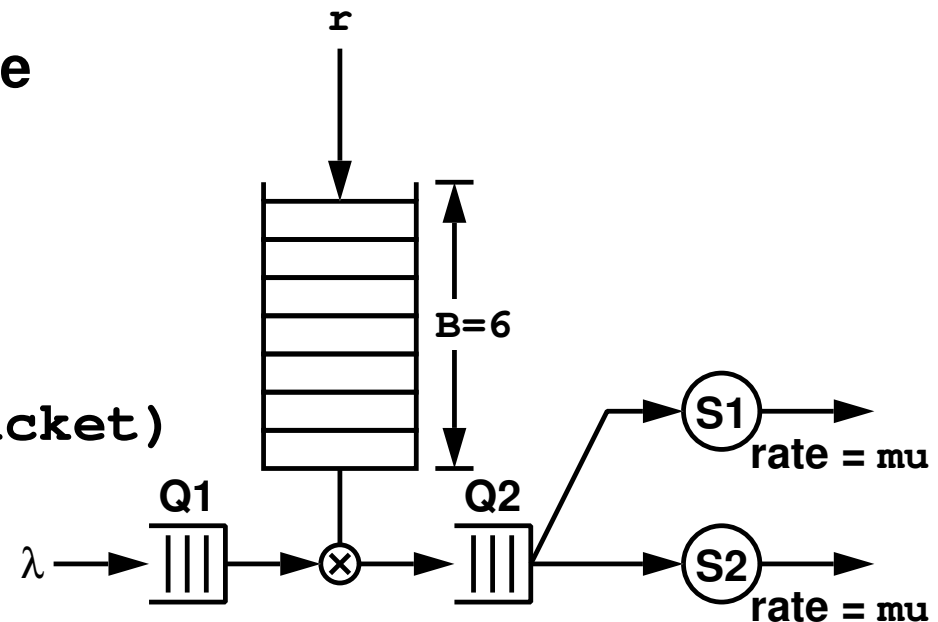
```
for (;;) {
    sleep
    generate a packet/token
    add packet/token to token bucket filter
}
```

— where must you lock and unlock mutex?

➡ Very high level pseudo-code for the *server thread*:

```
for (;;) {
    wait for packet in Q2
    remove packet from Q2
    sleep (to transmit packet)
}
```

— where must you lock and unlock mutex?



# Simulation/Emulation



## ➡ Packet thread pseudo-code (incomplete):

```

for (;;) {
    /* read a line from tsfile if in trace mode */
    get inter_arrival_time, tokens_needed, and service_time;
    /* calculate sleep time from inter_arrival_time */
    usleep(...);
    packet = NewPacket(tokens_needed, service_time, ...);
    pthread_mutex_lock(&mutex);
    Q1.enqueue(packet);
    ... /* other stuff */
    pthread_cond_broadcast(&cv);
    pthread_mutex_unlock(&mutex);
}

```

- ➡ must self-terminate as soon as this thread is no longer needed (i.e., no need to generate packets)
- ➡ must not call `pthread_cond_signal()`



# Simulation/Emulation



## ➡ Token thread pseudo-code (incomplete):

```

for (;;) {
    /* calculate sleep time from inter-token arrival time */
    usleep(...);
    pthread_mutex_lock(&mutex);
    tokens++;
    if (first packet in Q1 can now be moved into Q2) {
        packet = Q1.dequeue();
        Q2.enqueue(packet);
        pthread_cond_broadcast(&cv);
        tokens = 0; /* why? */
    }
    pthread_mutex_unlock(&mutex);
}

```

- ➡ must self-terminate as soon as this thread is no longer needed (i.e., no need to generate tokens)
- ➡ must not call `pthread_cond_signal()`







# Simulation/Emulation

- ➡ Server threads pseudo-code (incomplete):
- ➡ same first procedure for both server threads

```

for (;;) {
    /* wait for work */
    pthread_mutex_lock(&mutex);
    while (Q2.length() == 0 && !time_to_quit) {
        pthread_cond_wait(&cv, &mutex);
    }
    packet = Q2.dequeue();
    pthread_mutex_unlock(&mutex);
    /* work */
    usleep(packet.service_time);
}
  
```

- ➡ must self-terminate as soon as this thread is no longer needed (i.e., no need to transmit packets)




# Simulation/Emulation



**Many other requirements, for example:**

- = must move a packet at the correct time**
  - if a packet is eligible to be moved from Q1 to Q2, it must happen immediately
- = all threads must self-terminate when they are no longer needed**
- = drop packets**
  - if the token requirement for an arriving packet is too large (i.e.,  $> B$ ), must drop the packet
- = drop tokens**
  - if an arriving token finds a full bucket, the token is dropped
- = and many more...**
  - please read the spec yourself (don't get it from classmates)



# Program Printout

- ➡ Program output must look like what's in the spec
- you must **NOT** wait for emulation to end to print all these

## Emulation Parameters:

```

number to arrive = 20
lambda = 2          (if -t is not specified)
mu = 0.35           (if -t is not specified)
r = 4
B = 10
P = 3               (if -t is not specified)
tsfile = FILENAME   (if -t is specified)


```

← from commandline  
or tsfile

```

00000000.000ms: emulation begins
00000251.726ms: token t1 arrives, token bucket now has 1 token
00000502.031ms: token t2 arrives, token bucket now has 2 tokens
00000503.112ms: p1 arrives, needs 3 tokens, inter-arrival time = 503.112ms
00000503.376ms: p1 enters Q1
00000751.148ms: token t3 arrives, token bucket now has 3 tokens
00000751.186ms: p1 leaves Q1, time in Q1 = 247.810ms, token bucket now has 0 token
00000752.716ms: p1 enters Q2
00000752.932ms: p1 leaves Q2, time in Q2 = 0.216ms
00000752.982ms: p1 begins service at S1, requesting 2850ms of service
00001004.271ms: p2 arrives, needs 3 tokens, inter-arrival time = 501.159ms
00001004.526ms: p2 enters Q1
00001007.615ms: token t4 arrives, token bucket now has 1 token
00001251.259ms: token t5 arrives, token bucket now has 2 tokens
00001505.986ms: p3 arrives, needs 3 tokens, inter-arrival time = 501.715ms
00001506.713ms: p3 enters Q1
00001507.552ms: token t6 arrives, token bucket now has 3 tokens
00001508.281ms: p2 leaves Q1, time in Q1 = 503.755ms, token bucket now has 0 token
00001508.761ms: p2 enters Q2
00001508.874ms: p2 leaves Q2, time in Q2 = 0.113ms
00001508.895ms: p2 begins service at S2, requesting 1900ms of service
...

```



# Program Printout

```
...
00003427.557ms: p2 departs from S2, service time = 1918.662ms, time in system = 2423.286ms
00003612.843ms: p1 departs from S1, service time = 2859.861ms, time in system = 3109.731ms
...
?????????.???ms: p20 departs from S?, service time = ????.???ms, time in system = ????.???ms
?????????.???ms: emulation ends
```

## Statistics:


```
average packet inter-arrival time = <real-value>
average packet service time = <real-value>
```

```
average number of packets in Q1 = <real-value>
average number of packets in Q2 = <real-value>
average number of packets at S1 = <real-value>
average number of packets at S2 = <real-value>
```

```
average time a packet spent in system = <real-value>
standard deviation for time spent in system = <real-value>
```

```
token drop probability = <real-value>
packet drop probability = <real-value>
```

- ➡ **Timestamps** in the left column must have *microsecond* resolution
  - ➡ *measured time interval* must have *microsecond* resolution
  - ➡ use "%.6g" in printf() for <real-value>
- ➡ A value anywhere in the right column must be the exact differences between two corresponding timestamps



# Program Printout

```

...
00003427.557ms: p2 departs from S2, service time = 1918.662ms, time in system = 2423.286ms
00003612.843ms: p1 departs from S1, service time = 2859.861ms, time in system = 3109.731ms
...
?????????.???ms: p20 departs from S?, service time = ????.???ms, time in system = ????.???ms
?????????.???ms: emulation ends

```

## Statistics:

```

average packet inter-arrival time = <real-value>
average packet service time = <real-value>

```

```

average number of packets in Q1 = <real-value>
average number of packets in Q2 = <real-value>
average number of packets at S1 = <real-value>
average number of packets at S2 = <real-value>

```

```

average time a packet spent in system = <real-value>
standard deviation for time spent in system = <real-value>

```

```

token drop probability = <real-value>
packet drop probability = <real-value>


```

➡ Ex: why must the service time for p1 be exactly 2859.861ms?


$$3612.843\text{ms} - 752.982\text{ms} = 2859.861\text{ms}$$

— why so strict?


- your *printout* must be *self-consistent*



# Calculating Statistics




packet thread return from `usleep()`  
create packet (packet arrival)



**NOTE:** not all activities are shown


- ➡ A **packet** is not a thread, it's a **data structure**
  - ▬ it should have 7 timestamps to store "measured" information
  - ▬ it should also store "packet specification" (such as specified inter-arrival time, token requirement, service time)
    - these are **not** "measured" values
- ➡ Some packets need to be excluded from certain statistics
  - ▬ add to corresponding statistics only when a packet is **being elected**




# Mean and Standard Deviation

- ➡ **Average time**  
 = for  $n$  samples, add up all the time and divide by  $n$


- ➡ **Average number of packets at a server**  
 = same a fraction of time the server is busy



- ➡ **Average number of packets at Q1**




- ➡ **Standard deviation is the squareroot of variance**  
 =  $Var[X] = E[X^2] - (E[X])^2$   
 ○ must use the *population variance* equation



# Warmup #2 (Part 2)

Bill Cheng

*<http://merlot.usc.edu/william/usc/>*







# Handling <Ctrl+C> In Warmup #2





<Ctrl+C>

- packet arrival thread must stop generating packets and self-terminate
- token depositing thread must stop generating tokens and self-terminate
- a server thread must finish serving its current packet then self-terminate
- no packets or tokens must arrive
- need to take care of left-over packets that are stuck in Q1 and Q2
  - these are called "*removed packets*" because your simulator must remove them before you end your simulation
  - must distinguish 3 types of packets:
    - 1) *completed* packets
    - 2) *dropped* packets (packets that need  $> B$  tokens)
    - 3) *removed* packets



# Handling <Ctrl+C> In Warmup #2

- ➡ You can catch <Ctrl+C> asynchronously
- ▬ see lecture and understand why this is not a good idea
  - ▬ our recommendation is to *handle signals synchronously*
    - use a *signal-catching thread* and `sigwait()`
      - ◆ no signal handler at all
  - ▬ when <Ctrl+C> is caught, you should use the *pthread cancellation* mechanism to cancel packet and token threads
    - understanding thread cancellation will hopefully prevent you from make some mistakes in kernel assignment #1
      - ◆ although kernel cancellation is somewhat different from pthread cancellation



# Designate A Thread To Catch A Signal

- ➡ Look at the man pages of `pthread_sigmask()` on Ubuntu 16.04
- designate child thread to handler `SIGINT`
  - parent thread blocks `SIGINT` (some code deleted and use `SIGINT` instead of `SIGQUIT`):


```
#include <pthread.h>

int
main(int argc, char *argv[])
{
    pthread_t thread;
    sigset_t set;
    int s;

    sigemptyset(&set);
    sigaddset(&set, SIGINT);
    sigaddset(&set, SIGUSR1);
    s = pthread_sigmask(SIG_BLOCK, &set, NULL);
    if (s != 0)
        handle_error_en(s, "pthread_sigmask");

    s = pthread_create(&thread, NULL, &sig_thread, (void *) &set);
    if (s != 0)
        handle_error_en(s, "pthread_create");

    pause();
}
```



## pthread\_sigmask()



### Child thread example

- child thread unblocks SIGINT


```
static void *  
sig_thread(void *arg)  
{  
    sigset_t *set = arg;  
    int s, sig;  
  
    for (;;) {  
        s = sigwait(set, &sig);  
        if (s != 0)  
            handle_error_en(s, "sigwait");  
        printf("Signal handling thread got signal %d\n", sig);  
    }  
}
```

- child thread is designated to handle SIGINT and SIGUSR1, no other thread will get SIGINT and SIGUSR1



### Compile this code and run it

- press <Ctrl+C> to see that this program cannot be killed with <Ctrl+C>
- use "kill -15" command to kill this program



# Example In Lecture

```

some_state_t state;
sigset_t set;

main() {
    pthread_t thread;
    sigemptyset(&set);
    sigaddset(&set,
              SIGINT);
    sigprocmask(
        SIG_BLOCK,
        &set, 0);
    // main thread
    //      blocks SIGINT
    pthread_create(
        &thread, 0,
        monitor, 0);
    long_running_proc();
}

```


```

void long_running_proc() {
    while (a_long_time) {
        pthread_mutex_lock(&m);
        update_state(&state);
        pthread_mutex_unlock(&m);
        compute_more();
    }
}

void *monitor() {
    int sig;
    while (1) {
        sigwait(&set, &sig);
        pthread_mutex_lock(&m);
        display(&state);
        pthread_mutex_unlock(&m);
    }
    return(0);
}

```



= this is the **PREFERRED** way to catch SIGINT for warmup2



# sigwait()

```
int sigwait(sigset_t *set, int *sig)
```

- ➡ `sigwait()` blocks until a signal specified in `set` is received
  - return which signal caused it to return in `sig`
  - if you have a signal handler specified for `sig`, it will *not* get invoked when the signal is delivered
    - instead, `sigwait()` will return
- ➡ You should make sure that all the threads in your process have these signals blocked!
  - this way, when `sigwait()` is called, the calling thread temporarily becomes the *only* thread in the process who can receive the signal
- ➡ `sigwait(set)` *atomically unblocks* signals specified in `set` and *waits* for signal delivery




# Warmup2 Catching <Ctrl+C> Made Easy

- ➡ Use a *<Ctrl+C>-catching thread* and *block SIGINT everywhere else*
- when `sigwait()` returns because <Ctrl+C> has been pressed:

```
lock mutex
set global flag /* time to quit gracefully */
cancel packet and token threads
broadcast CV
unlock mutex
self terminate
```

- according to spec, you must not cancel server threads

- ➡ Remember, if you don't use a <Ctrl+C>-catching thread, your regular threads can be borrowed to deliver signals
- your code must deal with that and that can get very messy




# Warmup2 Cancellation Made Easy

- ➡ Only packet arrival and token depositing threads are allowed to be canceled
- use a <Ctrl+C>-catching thread (i.e., use `sigwait()`)
  - handling clean up with `pthread_cleanup_push()` and `pthread_cleanup_pop()` can be messy
  - it's good to use them in general, but for warmup2, maybe we should try to make our lives easier by observing this:
  - first procedure of packet or token thread looks like:

```
do-forever
    sleep();
    lock mutex
    /* stuff */
    unlock mutex
end-do
```

- if we have only one cancellation point, we know exactly where our threads will act on cancel






# Warmup2 Cancellation Made Easy

- what if we change it to:

```
disable cancellation
do-forever
  enable cancellation
  sleep();
  disable cancellation
  lock mutex
  /* stuff */
  unlock mutex
end-do
```

- where can this thread act on cancel?
- would this make our thread "*unresponsive*" when <Ctrl+C> is pressed?
- there is a race condition in this code
  - ◆ what if the SIGINT-catching thread cancels this thread right when this thread is waiting for mutex lock?
  - ◆ must not create a new packet/token once you have printed the "SIGINT caught" message




# Statistics - Running Average

- ➡ Do not store all transmitted packets in a list and calculate statistics at the end of your simulation
- it's a waste of memory
    - useful to learn how to write code with *small memory footprint*
  - keep a *running average* instead
    - just need one `int` and one `double`
      - ◆ `n` is the number of samples
      - ◆ `avg` is the average
      - ◆ when you get a new sample value, simply do:

```
avg = (n * avg + sample_value) / (n + 1)
n = n + 1
```

- avg above is  $E[X]$ , how do you calculate  $E[X^2]$ ?

```
avg2 = (n * avg2 + sample_value^2) / (n + 1)
n = n + 1
```



# Statistics - Running Average

— can you use running averages for all these statistics?

average packet inter-arrival time = <real-value>

average packet service time = <real-value>

average number of packets in Q1 = <real-value>

average number of packets in Q2 = <real-value>

average number of packets at S1 = <real-value>

average number of packets at S2 = <real-value>


average time a packet spent in system = <real-value>

standard deviation for time spent in system = <real-value>

token drop probability = <real-value>


packet drop probability = <real-value>

- what should correspond to a sample value?
- what types of packet must you consider?
  - ◆ avg packet inter-arrival time must consider **all** packets
  - ◆ other averages must only consider **completed** packets
  - ◆ packet drop probability is # dropped / # arrived




# **`gdb`**

- ➡ When you get a breakpoint in `gdb`, what happened with other threads?
- ➡ If you break inside a signal handler and you type "`where`", how come `gdb` seems confused?



# How To Learn New Concepts

- ➡ If there is a new concept that you are not familiar with, don't just try to write the final program
  - you won't know where the bugs are because you may not be clear about the concepts at multiple places
- ➡ Try writing small programs to test out ideas
  - try one idea at a time
  - use the debugger to get a better understanding of what's going on
  - then compile multiple ideas into one program and see if it works
- ➡ Ex:
  - `fork-wait.c`
  - `cat.c`
  - `redirect.c`
  - `thr-term.c`
  - `busywait.c`, `join.c`
  - `deadlock.c`, `trylock.c`
  - `whoopee.c`
  - `status-update.c`
  - `sigblock.c`, `sigwait.c`
  - `direct.c`
  - `cancel.c`
  - `three-threads.c`



# defs.h


```
#ifndef _DEFS_H_
#define _DEFS_H_

#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <signal.h>
#include <errno.h>

#include <pthread.h>


#ifndef NULL
#define NULL 0L
#endif /* ~NULL */

#endif /* _DEFS_H_ */
```



## fork-wait.c


```
#include "defs.h"
#define NUM_CHILD 5
int sleep_time[NUM_CHILD], chd_num=0;
int main(int argc, char *argv[])
{
    srand48(time(0));
    for (chd_num=0; chd_num < NUM_CHILD; chd_num++) {
        sleep_time[chd_num] = lrand48() % 5000000;
        if (fork() == 0) {
            int pid=((int)getpid());
            printf("(Child) pid = %1d (0x%08x)\n", pid, pid);
            usleep(sleep_time[chd_num]);
            exit(child_pid+1);
        }
    }
    for (;;) {
        int pid=0, rc=0;
        if ((pid=wait(&rc)) == (-1)) break;
        printf("child %1d exited: 0x%08x.\n", pid, rc);
    }
    return 0;
}
```



## cat.c

```
#include "defs.h"
#define BUFSIZE 1024
int main(int argc, char *argv[])
{
    char buf[BUFSIZE];
    int n=0;
    const char *note="Write failed\n";


    while ((n = read(0, buf, sizeof(buf))) > 0)
        if (write(1, buf, n) != n) {
            (void)write(2, note, strlen(note));
            exit(1);
        }
    return 0;
}
```






# redirect.c

```
#include "defs.h"
int main(int argc, char *argv[])
{
    pid_t pid=(pid_t)0;
    if ((pid=fork()) == 0) {
        close(1);
        if (open("/tmp/Output",
                O_CREAT|O_WRONLY,
                0666) == -1) {
            perror("/tmp/Output");
            exit(1);
        }
        execl("/bin/date", "date", (char*)0);
        exit(1);
    }
    while(pid != wait(0)) ;
    return 0;
}
```




# thr-term.c

```
#include "defs.h"
void *child(void *arg)
{
    if (*(int*)arg > 2) pthread_exit((void*)1);
    return((void*)2);
}
int main(int argc, char *argv[])
{
    pthread_t thread;
    void *result=NULL;
    pthread_create(&thread, 0, child, &argc);
    pthread_join(thread, (void**)&result);
    switch ((int)(long)result) {
    case 1: printf("result is 1\n"); break;
    case 2: printf("result is 2\n"); break;
    }
    return 0;
}
```



# busywait.c

```
#include "defs.h"
#define NUM_THRS 100
int done[NUM_THRS];
pthread_t tid[NUM_THRS];
void *child(void *arg)
{
    int index=(int)(long)(arg);
    usleep(lrand48()%10000000);
    done[index] = 1;
    return 0;
}
int main(int argc, char *argv[])
{
    int i=0;
    memset(done, 0, sizeof(done));
    srand48(0);
    for (i=0; i < NUM_THRS; i++)
        pthread_create(&tid[i], 0, child, (void*)(long)i);
    waitall();
    return 0;
}
```



# busywait.c

```
void waitall()
{
    for (;;) {
        int i=0, num_done=0;
        for (i=0; i < NUM_THRS; i++) {
            if (!done[i]) break;
            num_done++;
        }
        if (num_done == NUM_THRS) break;
    }
}
```



Why is this busy wait?


— the main thread is not doing anything useful



Fix?

— sleep for 100ms before checking

- to avoid doing busy wait
- not really a good solution




# join.c



Real fix

— join with all threads

```
void waitall()  
{  
    int i=0;  
    for (i=0; i < NUM_THRS; i++)  
        pthread_join(tid[i], 0);  
}
```




# deadlock.c

➡ Try to deadlock child thread and main thread

```
#include "defs.h"
void *child(void *arg)
{
    for (;;) { proc1(); }
    return (void*) 0;
}

int main(int argc, char *argv[])
{
    pthread_t thread;
    srand48(time(0));
    pthread_create(&thread, 0, child, 0);
    for (;;) { proc2(); }
    return 0;
}
```




# deadlock.c

```
pthread_mutex_t m1=PTHREAD_MUTEX_INITIALIZER;
pthread_mutex_t m2=PTHREAD_MUTEX_INITIALIZER;
```

```
void proc1() {
    pthread_mutex_lock(&m1);
    pthread_mutex_lock(&m2);
    printf("1");
    fflush(stdout);
    pthread_mutex_unlock(&m2);
    pthread_mutex_unlock(&m1);
    usleep(100000);
}
```

```
void proc2() {
    pthread_mutex_lock(&m2);
    pthread_mutex_lock(&m1);
    printf("2");
    fflush(stdout);
    pthread_mutex_unlock(&m1);
    pthread_mutex_unlock(&m2);
    usleep(100000);
}
```


- ➡ Why no deadlock?
  - threads are alternating
- ➡ How to make it deadlock?
  - call `printf("-")` after locking `m1` in `proc1()` and call `printf("+")` after locking `m2` in `proc2()`
    - deadlock right away! (why?)



# trylock.c

➡ How to use `trylock()` to avoid deadlock

```
void proc2() {  
    while (1) {  
        pthread_mutex_lock(&m2);  
        printf("+");  
        fflush(stdout);  
        if (!pthread_mutex_trylock(&m1))  
            break;  
        pthread_mutex_unlock(&m2);  
    }  
    printf("2");  
    fflush(stdout);  
    pthread_mutex_unlock(&m1);  
    pthread_mutex_unlock(&m2);  
    usleep(100000);  
}
```






# whoopee.c

➡ Let's catch SIGINT

```
#include "defs.h"
void handler(int signo)
{
    printf("Got signal %1d.  Whoopee!!\n", signo);
}
int main(int argc, char *argv[])
{
    sigset(SIGINT, handler);
    for (;;) { }
    return 1;
}
```



# status-update.c

```

#include "defs.h"
typedef struct foo {
    int x, y;
} my_state;
my_state state;
int main() {
    state.x = state.y = 0;
    sigset(SIGINT, handler);
    long_running_proc();
    return 0;
}

void long_running_proc() {
    int i=0;
    for (i=0; i < 100; i++) {
        update_state(&state);
        compute_more();
    }
}


void handler(int signo) {
    display(&state);
}

void update_state(
    my_state *ptr) {
    ptr->x++;
    usleep(100000);
    ptr->y++;
}

void display(my_state *ptr) {
    printf("x = %1d\n", ptr->x);
    usleep(1000);
    printf("y = %1d\n", ptr->y);
}

void compute_more() { usleep(100000); }

```




# sigblock.c

➡ Let's block SIGINT

```
#include "defs.h"
typedef struct foo {
    int x, y;
} my_state;
my_state state;
sigset_t set;
int main() {
    state.x = state.y = 0;
    sigemptyset(&set);
    sigaddset(&set, SIGINT);
    sigset(SIGINT, handler);
    long_running_proc();
    return 0;
}
```

```
void long_running_proc() {
    int i=0;
    for (i=0; i < 100; i++) {
        sigset_t old_set;
        sigprocmask(SIG_BLOCK,
            &set, &old_set);
        update_state(&state);
        sigprocmask(SIG_SETMASK,
            &old_set, 0);
        compute_more();
    }
}

void handler(int signo) {
    display(&state);
}
```



## sigwait.c

```

#include "defs.h"
typedef struct foo {
    int x, y;
} my_state;
my_state state;
sigset_t set;
pthread_mutex_t m=
PTHREAD_MUTEX_INITIALIZER;
int main() {
    pthread_t thr;
    state.x = state.y = 0;
    sigemptyset(&set);
    sigaddset(&set, SIGINT);
    sigprocmask(SIG_BLOCK,
        &set, 0);
    pthread_create(&thr, 0,
        monitor, 0);
    long_running_proc();
    return 0;
}


```

```

void long_running_proc() {
    int i=0;
    for (i=0; i < 100; i++) {
        pthread_mutex_lock(&m);
        update_state(&state);
        pthread_mutex_unlock(&m);
        compute_more();
    }
}

void *monitor(void *arg) {
    int sig=0;
    for (;;) {
        sigwait(&set, &sig);
        pthread_mutex_lock(&m);
        display(&state);
        pthread_mutex_unlock(&m);
    }
    return 0;
}

```



# direct.c



Direct a thread to catch SIGINT


— see "man pthread\_sigmask" on Ubuntu 16.04

```
#include "defs.h"
```

```
#define handle_error_en(en, msg) \
    do { errno = en; perror(msg); exit(1); } while (0)
```

```
static void *
sig_thread(void *arg)
{
    sigset_t *set = arg;
    int s, sig;

    for (;;) {
        s = sigwait(set, &sig);
        if (s != 0)
            handle_error_en(s, "sigwait");
        printf("Signal handling thread got signal %d\n",
            sig);
    }
}
```




# direct.c

```
int
main(int argc, char *argv[])
{
    pthread_t thread;
    sigset_t set;
    int s;

    sigemptyset(&set);
    sigaddset(&set, SIGINT);
    sigaddset(&set, SIGUSR1);
    s = pthread_sigmask(SIG_BLOCK, &set, NULL);
    if (s != 0)
        handle_error_en(s, "pthread_sigmask");

    s = pthread_create(&thread, NULL, &sig_thread,
        (void *) &set);
    if (s != 0)
        handle_error_en(s, "pthread_create");


    pause();
}
```



# cancel.c

## ➡ Cancellation


```
#include "defs.h"
#define NUM_THRS 10
pthread_t tid[NUM_THRS];
sigset_t set;
void cleanup(void *arg)
{
    int index=(int)(long)(arg);
    printf("Clean up thread %1d\n", index);
}
void *child(void *arg)
{
    pthread_cleanup_push(cleanup, arg);
    for (;;) {
        usleep(lrand48() % 1000000);
    }
    pthread_cleanup_pop(0);
    return 0;
}
```



# cancel.c

```
void waitall()
{
    int i=0;
    for (i=0; i < NUM_THRS; i++)
        pthread_join(tid[i], 0);
}


void *monitor(void *arg) {
    int i=0, sig=0;
    printf("Press <Ctrl+C>: ");
    fflush(stdout);
    sigwait(&set, &sig);
    printf("\nGot signal %1d\n", sig);
    for (i=0; i < NUM_THRS; i++) {
        while (tid[i] == 0) {
            usleep(100000);
        }
        pthread_cancel(tid[i]);
    }
    return 0;
}
```






## cancel.c

```
int main(int argc, char *argv[])
{
    pthread_t thr;
    int i=0;
    srand48(time(0));
    sigemptyset(&set);
    sigaddset(&set, SIGINT);
    sigprocmask(SIG_BLOCK, &set, 0);
    for (i=0; i < NUM_THRS; i++) {
        pthread_create(&tid[i], 0, child, (void*)(long)i);
    }
    pthread_create(&thr, 0, monitor, 0);
    waitall();
    pthread_join(thr, 0);
    return 0;
}
```




## three-threads.c

```
/*
 * Demonstrate that only a thread with SIGINT unblocked
 * will "see" SIGINT.
 * Currently, checking if num == 2 to unblock SIGINT
 * for thread 2. For all other threads, usleep() is
 * never interrupted.
 * Can change the comparison to 1 or 3 to demonstrate
 * how to designate a different thread to only
 * "see" SIGINT.
 */
#include <errno.h>
#include "defs.h"
sigset_t set;
void handler(int signo)
{
    printf("4");
    fflush(stdout);
}
```




## three-threads.c

```
void *child(void *arg) {
    int num=(int) (long) arg;
    if (num == 2) {
        sigprocmask(SIG_UNBLOCK, &set, 0);
    }
    for (;;) {
        printf("%1d", num);
        fflush(stdout);
        if (usleep(500000) != 0) {
            if (errno == EINTR) {
                printf("\nthread %1d interrupted...\n", num);
                fflush(stdout);
            }
        }
    }
    return 0;
}
```



# three-threads.c


```
int main(int argc, char *argv[])
{
    int i=0;
    pthread_t t[3];
    sigset(SIGINT, handler);
    sigemptyset(&set);
    sigaddset(&set, SIGINT);
    sigprocmask(SIG_BLOCK, &set, 0);
    for (i=0; i < 3; i++)
        pthread_create(&t[i], 0, child, (void*)(long)(i+1));
    for (;;) {
        printf("0");
        fflush(stdout);
        usleep(500000);
        if (errno == EINTR) {
            printf("\nMain thread interrupted...\n");
            fflush(stdout);
        }
    }
    return 0;
}
```



# Kernel Programming Assignments

Bill Cheng


*<http://merlot.usc.edu/william/usc/>*



# Kernel Programming Assignments (Part 1)

Bill Cheng

*<http://merlot.usc.edu/william/usc/>*




# Academic Integrity Policy




From the spec:

- ***accessing*** a submission from a previous semester is considered cheating
  - e.g., you run a submission from a previous semester to see how another student's code works
- if it's ***evident*** from your submission that you have ***accessed*** a submission from a previous semester, I will forward your case to USC Student Judicial Affairs for cheating investigation
- the standard punishment is a grade of **F** in the class.
- since there is only one submission from a team, ***even if only one team member cheated, the entire team is considered cheated.***
- therefore, please make sure that your teammates are not cheating!



# Academic Integrity Policy

- ➡ Since the kernel assignments are group assignments, if *one of your teammate cheated*, the *entire group* is considered cheated
- ➡ it is *your responsibility* to make sure that *your teammates are not cheating*
  - ➡ if you noticed that your teammate has written some code he/she cannot explain, you need to delete the code and write from scratch!
  - ➡ if you know that your teammate has code from previous semester, it's best you ask your teammate to *throw away the code*
- ➡ As mentioned in Lecture 1, posting `weenix` code (whether it's the original code or your code or a combination) to *public* bitbucket (or something similar) is considered *cheating*
- ➡ posting warmup code to public bitbucket is also not permitted
  - ➡ you must *not* use `github.com`
  - ➡ can use `bitbucket.org`, but make sure it's private





# Kernel Programming Assignments

- ➡ Tom Doeppner's *weenix* source and binary code
- provided as `weenix-assignment-3.8.0.tar.gz`
  - incomplete
  - contains code like:

```
NOT_YET_IMPLEMENTED("PROCS: bootstrap");
```


*assignment name* ————↑

*function name* —————↑

- your job is to implement these functions by *replacing* these lines with your code
  - ◆ please *replace* them *in-place*
  - ◆ do NOT *reorder* code there

- to look for such code:

```
grep PROCS: kernel/*.c
grep PROCS: kernel/*/*.c
grep PROCS: kernel/*/*/*.c
grep PROCS: kernel/*/*/*/*.c
```




# Download and Setup

```
% tar xvzf weenix-assignment-3.8.0.tar.gz
% cd weenix-assignment-3.8.0/weenix
% make clean
% make
% ./weenix -n
```

- ➡ Don't type the long commands, just copy from your web browser and paste into your terminal
- ▬ if something looks wrong, seek help as soon as you can
  - ▬ best *not* to do the above in a Shared Folder
    - see kernel FAQ if you have to do that

- ➡ Don't forget to install these when you setup Ubuntu 16.04:

```
sudo apt-get install -y git
sudo apt-get install -y gitk
sudo apt-get install -y git-core
sudo apt-get install -y qemu
sudo apt-get install -y cscope
```




# Download and Setup

- ➡ If all goes well, you should see tons of stuff fly by and the following at the bottom of the terminal:

```
Not yet implemented: PROCS: bootstrap, file
main/kmain.c, line 184
panic in main/kmain.c:186 bootstrap():
weenix returned to bootstrap()!!! BAD!!!
```


Kernel Halting.

- ➡ = press <Ctrl+C> in the terminal to quit
- ➡ Make sure you have tried the above this weekend
  - ➡ = any problem, let us know **NOW**
- ➡ Don't worry that you don't know how to do kernel 1
  - ➡ = for now, just get familiar with the spec and the documentation
    - and may be read some kernel code
    - you are **not** expected to understand a lot of the kernel code



# Documentation


- ➡ The `weenix` documentation is in `doc/latex/documentation.pdf`
  - ▬ introduces `weenix` to you
  - ▬ detailed instructions on assignments
  - ▬ you must read it thoroughly
- ➡ We are doing three of the assignments
  - ▬ *Processes and Threads (PROCS)*
  - ▬ *Virtual File System (VFS)*
  - ▬ *Virtual Memory (VM)*
- ➡ We are *not* doing two of the assignments
  - ▬ Drivers (DRIVERS)
  - ▬ System V File System (S5FS)
  - ▬ these are done for you and they are compiled and provided as libraries
    - `kernel/libdrivers.a` and `kernel/libS5fs.a`
      - ◆ source code for these are *not* available; you need to learn how to *work around* code you don't have



# Kernel Programming Assignments (Part 2)

Bill Cheng

*<http://merlot.usc.edu/william/usc/>*



# Compilation and Configuration

➡ **Config.mk** controls what gets compiled and configured into the kernel (weenix is a monolithic OS)

- ▬ for PROCS, use the original **Config.mk**
  - set **DRIVERS** to 1 to complete this assignment
- ▬ for VFS, set **DRIVERS** and **VFS** to 1
- ▬ for VM, set **DRIVERS**, **VFS**, **S5FS**, and **VM** to 1
  - set **VM** to 0 at first to get `kernel/mm/pframe.c` working
  - then set **VM** to 1 to work on the rest of the assignment
  - set **DYNAMIC** to 1 in the end if everything is working
- ▬ by default: **DBG = all**
  - the grader will use these for *grading*:

**DBG = error, test**


**DBG = error, print, test**

➡ Every time you modify **Config.mk**, you should do:

**% make clean**

**% make**

**% ./weenix -n**




# Debugging with gdb

- ➡ The `weenix` documentation says to do this to debug the kernel:
- ```
% ./weenix -n -d gdb
```
- although you can use `gdb`, but you cannot see kernel debugging messages (from `dbg()` calls)
- ➡ To see kernel debugging messages AND debug the kernel, do:
- set `GDBWAIT=1` in `Config.mk` then **recompile** kernel
  - run `weenix` under `gdb` with:

```
% ./weenix -n -d gdb -w 10
```

    - if you have a slow machine, you should use a larger value
    - if you have a fast machine, you should use a smaller value
  - unfortunately, if you have compiled with `GDBWAIT=1` and want to run **without `gdb`**, `weenix` will **freeze**
    - you have to set `GDBWAIT` back to 0 and recompile if you want to run `weenix` without `gdb`
    - you should set `GDBWAIT` back to 0 when you submit your assignment for grading




# Submissions



## Processes and Threads (PROCS)

```
% make procs-submit  
tar cvzf procs-submit.tar.gz \  
    Config.mk \  
    procs-README.txt \  
    ...
```


- must fill out `procs-README.txt`, it's your assignment's documentation
  - **IMPORTANT:** you must copy the `k1-README.txt` template in the spec into `procs-README.txt` and make changes
  - this is where you should also include
    - 1) how to split the points (in terms of percentages and must sum to 100%)
    - 2) brief justification about the split (if not equal split)
      - ◆ please understand that if I have to get involved, it can also be unfair since I won't have all the information
      - ◆ best to claim even splits
- submit `procs-submit.tar.gz` using web form





# Submissions

- ➡ Virtual File System (VFS)
  - need to fill out `vfs-README.txt` (start with `k2-README.txt`)
    - start with the `k2-README.txt` template in the spec
  - `% make vfs-submit`
- ➡ Virtual Memory (VM)
  - need to fill out `vm-README.txt` (start with `k3-README.txt`)
    - start with the `k3-README.txt` template in the spec
  - `% make vm-submit`
- ➡ Must **NOT** include **ANY OTHER** file not mentioned in "make" above
  - or we will **delete** it before grading
  - you must not add any files
- ➡ Submit source code only
  - we will deduct 2 points if you submit binary files
  - we will deduct 2 points if you submit extra files
  - we will deduct 2 points if you do not keep the same directory structure




# Verify Your Kernel Submission

- ➡ Assume that in your home directory, you have
- a *pristine* weenix-assignment-3.8.0.tar.gz
  - your submission file, e.g., procs-submit.tar.gz


- ➡ Do the following to verify your submission

```
% rm -rf /tmp/xyzzzy
% mkdir /tmp/xyzzzy
% cd /tmp/xyzzzy
% tar xvzf ~/weenix-assignment-3.8.0.tar.gz
% cd weenix-assignment-3.8.0/weenix
% tar xvzf ~/procs-submit.tar.gz
% make clean
% make
% ./weenix -n
[ go through grading guidelines line by line ]
[ check every line of your README file ]
```




# Grading

- ➡ The weenix code and comments and the weenix documentation are from Brown University
  - these are just hints and not our "requirements"
- ➡ Just like the warmup assignments, the *spec* and the *grading guidelines* are the requirements for our CS 402
  - they are related to "grading"
- ➡ Read the grading guidelines carefully!
  - e.g., altering or removing top comment block in a submitted `.c` file will **cost you 20 points** each!
    - why such stiff penalty?
      - ◆ because it's extremely easy to follow this rule
  - e.g., altering/removing a call to `dbgq()` in `bootstrap()` in your submitted `kmain.c` file will **cost you 20 points**
    - this is a signature showing who downloaded the kernel source
  - designate a teammate to check your submission against **every line in the grading guidelines**



# Structure Of Grading Guidelines

- ➡ "Plus Points" section of the grading guidelines
  - ➡ *mandatory KASSERTs*: section (A)
    - we are actually trying to help you to write some of your kernel code and give you some easy points at the same time!
    - add `KASSERT ()` calls (followed by a "**conforming `dbg ()` call**")
  - ➡ *SELF-checks*: last section
    - *every code sequence* inside any function you wrote to replace a `NOT_YET_IMPLEMENTED ()` call must **END** with a "**conforming `dbg ()` call**"
      - ◆ must use "*static analysis*" (independent of how your code will actually execute) to analyze your code
  - ➡ *PRE-CANNED tests*: other sections
- ➡ Please read the grading guidelines very carefully
  - ➡ when in doubt, ask the instructor!




# "Conforming dbg ( ) Call"

## ➡ "Conforming dbg ( ) call"

— general form:

```
dbg (DBG_PRINT, " (GRADING#S X.Y) \n" ) ;
```

- # is the kernel assignment number, can only be 1, 2, or 3
- s is the section number of the grading guidelines
  - for section (A) of the grading guidelines, you must use the corresponding x and y
    - ◆ read the requirements very very carefully!
  - for other sections of the grading guidelines:
    - ◆ x is a subtest number (applicable only when the subtest can be run separately with a shell-level command)
    - ◆ never use x.y
- you must *format it exactly* according to spec or you will lose a lot of points
  - read the requirements very very carefully!
- when in doubt, ask the instructor!



# "Conforming dbg ( ) Call"

- ➡ There are only **two** reasons to make a "conforming dbg ( ) call"
- 1) get credit for an item in **section (A)** of the grading guidelines
    - ◆ must use this form (must specify #, x, and y):

```
dbg (DBG_PRINT, " (GRADING#A X.Y) \n" );
```


- 2) get credit for **SELF-checks**

- ◆ must use this form (use x if applicable):

```
dbg (DBG_PRINT, " (GRADING#A) \n" );
dbg (DBG_PRINT, " (GRADING#B X) \n" );
dbg (DBG_PRINT, " (GRADING#C X) \n" );
...
```


- ◆ dbg (DBG\_PRINT, " (GRADING#A) \n" ); means start and stop the kernel (without doing anything else)

— when in doubt, more is better than less!



# "SELF-checks"

- ➡ As part of our requirements, you must *not* put/leave useless stuff in your kernel (i.e., don't leave trash in the kernel)
- *every code sequence must be traversable*
    - a code sequence is a block of code that does not contain conditionals or gotos and has only one entry point at the top
      - ◆ a code path is a path your code *may* take, i.e., there is a way to execute your code along that path
    - if a code sequence is not traversable, you must *delete* it
      - ◆ *must not leave useless code in the kernel!*
    - if you cannot demonstrate that there is way to get to it, you must *remove* it or we will take points off
- ➡ This is referred to as "SELF-checks" in the spec
- when you are confused about "SELF-checks", please come back and read this slide again to remind yourself why we are doing "SELF-checks" (and hopefully, that will answer your questions)



# "SELF-checks"


## ➡ What's useless code?

```
if (cond) {
    /* code seq */
}
```

- if cond can never be true, then code seq is *useless*
- to demonstrate that code seq is useful, you must tell the grader *which test to run* to reach the *END* of code seq
  - must add "*conforming dbg () call*" at the *END* of code seq

```
if (cond) {
    /* code seq */
    dbg(DBG_PRINT, " (GRADING#S X) \n");
}
```

- ◆ # is assignment number: 1, 2, or 3
- ◆ s is section number: B, C, D, ...
- ◆ x is subsection number (if applicable)
- just need to tell the grader *one way* to get there






# Must Use Static Analysis For SELF-checks

- ➡ To determine where you must add "conforming dbg() call", you must perform "**static analysis**" of your code (i.e., does not depend on how your code actually runs)

```
while (cond) {  
    /* seq1 */  
}  
/* seq2 */
```

- you may argue that the first time the while loop is executed, `cond` will be true, then later on, `cond` will be false, so you just need a "conforming dbg() call" at the end of `seq2`
  - that would be an incorrect analysis because it depends on how your code would execute
  - in this case, you must put a "conforming dbg() call" at the END of `seq1` and at the END of `seq2`

➡ Please read the spec for details!



# "Conforming dbg () Call" Requirement

➡ Example from spec:


```
/* sequence1 */
if (cond1) {
    /* sequence2 */
} else {
    /* sequence3 */
}
/* sequence4 */
```



```
/* sequence1 */
if (cond1) {
    /* sequence2 */
    /* conforming dbg() call */
} else {
    /* sequence3 */
    /* conforming dbg() call */
}
/* sequence4 */
/* conforming dbg() call */
```


➡ A "sequence" is a list of C statements *not* containing conditionals, gotos, or a label which is the target of a goto statement

- ➡ if you are not sure, you can make a "conforming dbg () call" at the **END** of **every code sequence**
  - you cannot go wrong with more!




# Backing Up Your Work & Collaboration

- ➡ You *have to* have a plan to *backup your code* and backup routinely
  - ▬ if you lose your work, no one can recover your files
  - ▬ can use *private* DropBox / iCloud / Microsoft cloud
  - ▬ or use a *private* bitbucket (must *not* use github)
    - share it among your team members only
- ➡ One simple way to backup your work
  - ▬ at the end of each day, do:
    - % make backup
    - it will tell you the name of the backup file
    - if you have a "Shared-ubuntu" shared folder in your home directory, your backup file will be copied there
      - ◇ if not, your backup file will be in the current directory and you should copy it into a shared folder
  - ▬ use your host's cloud backup facility to back up this file/folder



# Backing Up Your Work & Collaboration


- ➡ You should use `git` to collaborate among project partners
  - read "Pro Git" (a free online book, one of our textbooks)
- ➡ But you need a shared repository in the cloud to collaborate with your teammates
  - there are free `git` repositories on the web
    - unfortunately, most of the free ones are required to be visible by the world - you must **not** use these
      - ◆ on `github.com`, private repository automatically becomes public after 2 years (if you don't pay)
      - ◆ this is why you must **not** use `github.com`
    - you can use `bitbucket.org` but you need to make sure that your files remain private
  - apply for **free academic account** on `bitbucket.org` to share with teammates
    - make sure your projects are truly **private**



# Backing Up Your Work & Collaboration


➡ If you have two people working on the same file and then update the repository one after another

- ➡ `git` will attempt to merge the changes, but it may not be what you want
- ➡ may be it's best to *coordinate* and not have two people modifying the same file




# Early and Late Policies

- ➡ Same early submission policy as warmup projects
- ➡ Similar late submission policy as warmup projects
  - except that kernel 1 can be submitted by the kernel 2 deadline and get a 50% deduction
  - similarly, kernel 2 can be submitted by the kernel 3 deadline and get a 50% deduction
  - kernel 3 has regular late submission policy




# Extra Credit

- ➡ You can get *extra credit* for posting *timely* and *good/useful* answers to the class Google Group in response to questions posted by other students regarding *kernel* programming assignments
- ➡ the maximum number of extra credit points you can get is **10** points for each of the kernel assignments (on a 100-point scale)
  - ➡ "timely" means *within 8 hours of the original post*
    - if you don't "*reply*" to another student's post, you are not eligible for this type of extra credit
  - ➡ if you just repeat what others are saying, you are being "helpful" but what you post will not be "useful" since it's already been said



# How Do You Start?


- ➡ Definitely start with the documentation, spec, and kernel FAQ
- ➡ Read code, read lots and lots of code
  - try things out and see what happens (debugging statements)
  - you need to *absorb* other people's code, make sense of it
    - although you *don't have to understand everything*
  - that's what OS hacking (in the good sense) is all about
    - it's *not* about "implementing an OS"
- ➡ It's the *process* that matters
  - *not* the *answers*
  - it's about learning how to figure out *which* 2 or 3 lines (or 20 or 30 lines) of code to insert and *where*
- ➡ So, it needs to be experienced
  - you should not expect quick/straight answers
  - this is not an OS hacking class
- ➡ Learning to write OS code is like...






# How Do You Start?

- ➡ Definitely start with the documentation, spec, and kernel FAQ
- ➡ Read code, read lots and lots of code
  - try things out and see what happens (debugging statements)
  - you need to *absorb* other people's code, make sense of it
    - although you *don't have to understand everything*
  - that's what OS hacking (in the good sense) is all about
    - it's *not* about "implementing an OS"
- ➡ It's the *process* that matters
  - *not* the *answers*
  - it's about learning how to figure out *which* 2 or 3 lines (or 20 or 30 lines) of code to insert and *where*
- ➡ So, it needs to be experienced
  - you should not expect quick/straight answers
  - this is not an OS hacking class
- ➡ Learning to write OS code is like... *Zen*




# Getting Help

- ➡ If you have questions about the kernel assignments
- read documentation, textbook, lecture slides, read more code
  - send me e-mail
    - please understand that neither I nor other teaching staff can *tell you what code to write!*
    - if you ask me if you should set this variable to x or y, I will ask you to try both and figure out what makes more sense!
      - ◆ if you send us questions like that, we may simply forward your post to the class Google Group since we cannot tell you what code to write
  - post your questions to the class Google Group
    - your classmates are a great resource!
    - sometimes, we may not immediately answer these questions to give your classmates an opportunity to earn extra credit points
      - ◆ we may wait 2 hours



# Pitfalls To Avoid

- ➡ Your team need to *meet often*
  - once a day is preferred
    - work at the same place at the same time
    - have lots of discussions (and write a fair amount of code)
  - swallow your pride, be honest with your teammates, don't hide your weakness
    - everyone gets the same grade
    - if no one is really good at this (which is not unusual), someone (or more) has to step up
- ➡ You don't have to know what every piece of code is doing
  - learn how to *assume* that other code works (until proven otherwise)
    - other code works kind of like what's covered in lectures
  - use "grep" to *get an idea* of how a function is used and how a field in a data structure is used



# Kernel 1


Bill Cheng

*<http://merlot.usc.edu/william/usc/>*



**Prerequisite: a simple system (Ch 4)**

**— when you finish the prerequisite, please come back and review this material**




# Boot Prestine Kernel

```
% script
Script started, file is typescript
reading .login (xterm) ...
% ./weenix -n
/usr/bin/qemu-system-i386
...
Not yet implemented: PROCS: bootstrap, file \
main/kmain.c, line 184
panic in main/kmain.c:186 bootstrap(): weenix \
returned to bootstrap()!!! BAD!!!
```


Kernel Halting.

```
^C
qemu: terminating on signal 2
% exit
Script done, file is typescript
% more typescript
```



# Look At "typescript"

- ➡ Where is kernel's text segment?
  - ▬ `0xc0000000-0xc0038000` means `[0xc0000000, 0xc0038000)`
    - this interval is "closed" on the left (i.e., includes) and "open" on the right (i.e., excludes)
      - ◆ pretty much all intervals are denoted this way
- ➡ Where are the other segments?
  - ▬ data: `[0xc0038000-0xc0044baa)`
  - ▬ bss: `[0xc0044baa-0xc005a000)`
  - ▬ page system: `[0xc00a0000-0xcfe9d000)`
    - I think this is like the dynamic segment for the kernel (i.e., Buddy System)
- ➡ If your numbers are slightly different, it's probably okay



# Look At "typescript"




What else?

- kernel data structures are initialized (e.g., memory allocators)
- some hardware are initialized (e.g., PIC)
  - *don't need to understand EVERYTHING*
- if your kernel access anywhere outside this range, most likely you will get a **kernel page fault**
  - this will be followed by a **kernel panic**



**VERY IMPORTANT:** read kernel FAQ about *"how can I debug a page fault that caused kernel panic?"*


- whenever you get a kernel page fault, you must follow the steps there to figure out **EXACTLY** where your kernel crashed
  - I have no special power, I have no idea why your kernel crashed



# Kernel 1

% make nyi

|                        |                              |       |
|------------------------|------------------------------|-------|
| proc/kmutex.c:36       | kmutex_init()                | PROCS |
| proc/kmutex.c:48       | kmutex_lock()                | PROCS |
| proc/kmutex.c:58       | kmutex_lock_cancellable()    | PROCS |
| proc/kmutex.c:78       | kmutex_unlock()              | PROCS |
| proc/kthread.c:106     | kthread_create()             | PROCS |
| proc/kthread.c:127     | kthread_cancel()             | PROCS |
| proc/kthread.c:148     | kthread_exit()               | PROCS |
| proc/proc.c:223        | proc_create()                | PROCS |
| proc/proc.c:254        | proc_cleanup()               | PROCS |
| proc/proc.c:268        | proc_kill()                  | PROCS |
| proc/proc.c:280        | proc_kill_all()              | PROCS |
| proc/proc.c:294        | proc_thread_exited()         | PROCS |
| proc/proc.c:315        | do_waitpid()                 | PROCS |
| proc/proc.c:328        | do_exit()                    | PROCS |
| proc/sched.c:121       | sched_cancellable_sleep_on() | PROCS |
| proc/sched.c:137       | sched_cancel()               | PROCS |
| proc/sched.c:179       | sched_switch()               | PROCS |
| proc/sched.c:198       | sched_make_runnable()        | PROCS |
| proc/sched_helper.c:43 | sched_sleep_on()             | PROCS |
| proc/sched_helper.c:49 | sched_wakeup_on()            | PROCS |
| proc/sched_helper.c:56 | sched_broadcast_on()         | PROCS |
| main/kmain.c:184       | bootstrap()                  | PROCS |
| main/kmain.c:279       | initproc_create()            | PROCS |
| main/kmain.c:298       | initproc_run()               | PROCS |





# Kernel 1

## ➡ Kernel thread creation, cancellation, and destruction

|                    |                  |       |
|--------------------|------------------|-------|
| proc/kthread.c:106 | kthread_create() | PROCS |
| proc/kthread.c:127 | kthread_cancel() | PROCS |
| proc/kthread.c:148 | kthread_exit()   | PROCS |


— please keep MPT=0 in Config.mk and **only** implement *single thread processes*

## ➡ Kernel scheduler

|                        |                              |       |
|------------------------|------------------------------|-------|
| proc/sched.c:121       | sched_cancellable_sleep_on() | PROCS |
| proc/sched.c:137       | sched_cancel()               | PROCS |
| proc/sched.c:179       | sched_switch()               | PROCS |
| proc/sched.c:198       | sched_make_runnable()        | PROCS |
| proc/sched_helper.c:43 | sched_sleep_on()             | PROCS |
| proc/sched_helper.c:49 | sched_wakeup_on()            | PROCS |
| proc/sched_helper.c:56 | sched_broadcast_on()         | PROCS |

## ➡ Kernel process creation, waiting, and destruction

|                 |                      |       |
|-----------------|----------------------|-------|
| proc/proc.c:223 | proc_create()        | PROCS |
| proc/proc.c:254 | proc_cleanup()       | PROCS |
| proc/proc.c:268 | proc_kill()          | PROCS |
| proc/proc.c:280 | proc_kill_all()      | PROCS |
| proc/proc.c:294 | proc_thread_exited() | PROCS |
| proc/proc.c:315 | do_waitpid()         | PROCS |
| proc/proc.c:328 | do_exit()            | PROCS |



# Kernel 1

## ➡ Kernel mutex

|                  |                           |       |
|------------------|---------------------------|-------|
| proc/kmutex.c:36 | kmutex_init()             | PROCS |
| proc/kmutex.c:48 | kmutex_lock()             | PROCS |
| proc/kmutex.c:58 | kmutex_lock_cancellable() | PROCS |
| proc/kmutex.c:78 | kmutex_unlock()           | PROCS |


## ➡ Kernel startup

|                  |                   |       |
|------------------|-------------------|-------|
| main/kmain.c:184 | bootstrap()       | PROCS |
| main/kmain.c:279 | initproc_create() | PROCS |
| main/kmain.c:298 | initproc_run()    | PROCS |

➡ Read the *comment blocks* to figure out what these functions suppose to do and how they are related to each other

- use "grep" to see how they are called
- feel free to discuss this in the class Google Group
- you can talk about code that came with the prestine kernel
  - do *not* post code/pseudo-code you are planning to write

➡ When a thread gives up the CPU, you must make sure that all *global variables* and *data structures* are in a *consisten state*




# How To Start?



I would recommend doing this first:

- phase 1: get the kernel to simply start and quit with `DRIVERS=0`
  - in `bootstrap()`, create IDLE proc and switch to IDLE proc
    - ◆ first procedure of IDLE proc is written for you already, don't change anything
    - ◆ *Hint*: look at the code in `faber_thread_test()` to see how to create a process, create a thread in it, run it
    - ◆ Note: code in `faber_thread_test()` is running in *thread context*, you are *not* in a thread context in `bootstrap()`
  - in `initproc_create()`, create INIT proc and INIT thread
  - don't write code in `initproc_run()`, which is the first procedure of INIT proc
    - ◆ it should self-terminate
    - ◆ single-step to see where it goes
  - this should wake up the IDLE proc, which will turn off the machine




# How To Start?

➡ I would recommend doing this first:

- phase 1: get the kernel to simply start and quit with DRIVERS=0
  - make sure the *kernel halts cleanly*
  - these functions are involved (the list may *not* be complete):

|                        |                              |       |
|------------------------|------------------------------|-------|
| proc/kthread.c:106     | kthread_create()             | PROCS |
| proc/kthread.c:148     | kthread_exit()               | PROCS |
| proc/proc.c:223        | proc_create()                | PROCS |
| proc/proc.c:254        | proc_cleanup()               | PROCS |
| proc/proc.c:294        | proc_thread_exited()         | PROCS |
| proc/proc.c:315        | do_waitpid()                 | PROCS |
| proc/sched.c:121       | sched_cancellable_sleep_on() | PROCS |
| proc/sched.c:179       | sched_switch()               | PROCS |
| proc/sched.c:198       | sched_make_runnable()        | PROCS |
| proc/sched_helper.c:43 | sched_sleep_on()             | PROCS |
| proc/sched_helper.c:49 | sched_wakeup_on()            | PROCS |
| proc/sched_helper.c:56 | sched_broadcast_on()         | PROCS |
| main/kmain.c:184       | bootstrap()                  | PROCS |
| main/kmain.c:279       | initproc_create()            | PROCS |

- yes, it's a lot of code to get working just for phase 1!
- in a way, phase 1 is the most important step
  - ◆ if you do it the wrong way, you will have to come back and fix your code




# How To Start?



I would recommend doing this first:

- phase 2: get the kernel to simply start and quit *cleanly*
  - call `faber_thread_test()` from `initproc_run()`
    - ◆ start with `CS402TESTS=1` in `Config.mk`
    - ◆ then set `CS402TESTS=2, 3`, and so on, up to 8
  - always make sure the *kernel halts cleanly*
- phase 3: set `DRIVERS=1` in `Config.mk`
  - run `kshell` in `initproc_run()`
    - ◆ "help", "echo" and "exit" `kshell` commands should work
  - add `kshell` commands to invoke any test function in grading guidelines and your README (see rules about "SELF-checks")
    - ◆ for each `kshell` command, you need to *create a child process* and set the test function as the *first procedure* of the thread in the child process




# How To Start?

➡ Please remember that the teaching staff *cannot tell you what code to write*

➡ Here's what's appropriate to talk about in class Google Group

- 1) the spec
- 2) the kernel FAQ
- 3) the grading guidelines
- 4) test code

- kernel 1: `faber_thread_test()`, `sunghan_test()`, `sunghan_deadlock_test()`
  - ◆ they are mentioned in the grading guidelines
  - ◆ your kernel needs to work with these test code perfectly



# Hints?




Hints are all over the place!

- documentation
- spec
- comment in code
- kernel code itself
- kernel FAQ



For example, the spec says:

- the kernel is very very powerful
  - if there is a bug, it's *your* bug!
- the weenix kernel is *non-preemptive*
  - *non-preemptive* means that a thread *cannot be preempted* by *another thread*
    - ◆ it can be interrupted to *service an interrupt*, then goes back to what it was doing before
    - ◆ if a kernel thread *does not want to be cancelled*, there is *no way* to kill it
- we are *not* implementing multiple threads per process, i.e., MTP=0 in Config.mk



# "kmain.c"

## kmain()

```
context_setup(&bootstrap_context, bootstrap, 0, NULL, bstack,  
             PAGE_SIZE, bpdire);  
context_make_active(&bootstrap_context);  
panic("\nReturned to kmain()!!!\n");
```

## bootstrap()

```
NOT_YET_IMPLEMENTED("PROCS: bootstrap");  
panic("weenix returned to bootstrap()!!! BAD!!!\n");
```

## idleproc\_run()


```
kthread_t *initthr = initproc_create();  
init_call_all();  
GDB_CALL_HOOK(initialized);  
intr_enable();  
sched_make_runnable(initthr);  
child = do_waitpid(-1, 0, &status);
```

## initproc\_create()

```
NOT_YET_IMPLEMENTED("PROCS: initproc_create");
```

## initproc\_run()

```
NOT_YET_IMPLEMENTED("PROCS: initproc_run");
```






# More On Kernel 1


Bill Cheng

*<http://merlot.usc.edu/william/usc/>*



# Kernel Code

- ➡ We will go over *some* kernel code now (we will not put all the code in "faber\_test.c" and "sunghan\_test.c" on these slides)
- we will not cover all the test code
    - we will probably cover a *small number of test cases*
    - you need to learn how to read all these code and use them to figure out what you need to do
  - if you have a specific test you want the instructor to talk about next Friday, please send him an e-mail before 5pm on Thursday
- ➡ **IMPORTANT:** at any line in `faber_thread_test()`, ask yourself
- 1) where are all the threads/processes?
    - ◆ i.e., in which *queue* is a thread sleeping
  - 2) if a thread is not in the run queue, *who* is going to *unblock* it and *when/how* (by calling what function)?
  - 3) exactly how and where will a particular process/thread die?




## faber\_thread\_test()

- ➡ Let's look at the first set of tests in `faber_thread_test()`
- ▬ first subtest

```
dbg(DBG_TEST, "waitpid any test\n");
start_proc(&pt, "waitpid any test", waitpid_test, 23);
wait_for_any();

/*
 * Create a process and a thread with the given name
 * and calling teh given function. Arg1 is passed to
 * the function (arg2 is always NULL). The thread
 * is immediately placed on the run queue. A
 * proc_thread_t is returned, giving the caller a
 * pointer to the new process and thread to
 * coordinate tests. NB, the proc_thread_t is
 * returned by value, so there are no stack problems.
 */
static void start_proc(proc_thread_t *ppt, char *name,
    kthread_func_t f, int arg1) {
    ...
}
```




## faber\_thread\_test()

- ➡ Let's look at the first set of tests in `faber_thread_test()`
- ▬ first subtest

```
static void start_proc(proc_thread_t *ppt, char *name,
    kthread_func_t f, int arg1) {
    proc_thread_t pt;

    pt.p = proc_create(name);
    pt.t = kthread_create(pt.p, f, arg1, NULL);
    KASSERT(pt.p && pt.t &&
        "Cannot create thread or process");
    sched_make_runnable(pt.t)
    if (ppt != NULL) {
        memcpy(ppt, &pt, sizeof(proc_thread_t));
    }
}


void *waitpid_test(int arg1, void *arg2) {
    do_exit(arg1);
    return NULL;
}
```



## faber\_thread\_test()

- ➡ Let's look at the first set of tests in `faber_thread_test()`
- ▬ first subtest

```
/**
 * Call waitpid with a -1 pid and print a message
 * about any process that exits.
 * Returns the pid found, including -ECHILD when this
 * process has no children.
 */
static pid_t wait_for_any() {
    int rv;
    pid_t pid = do_waitpid(-1, 0, &rv);
    if (pid != -ECHILD)
        dbg(DBG_TEST, "child (%d) exited: %d\n", pid, rv);
    return pid;
}
```




## faber\_thread\_test()

- ➡ Let's look at the first set of tests in `faber_thread_test()`
- ▬ 2nd subtest

```
dbg(DBG_TEST, "waitpid test\n");
start_proc(&pt, "waitpid test", waitpid_test, 32);
pid = do_waitpid(2323, 0, &rv);
if ( pid != -ECHILD )
    dbg(DBG_TEST, "Allowed wait on non-existent pid\n");
wait_for_proc(pt.p);
```

```
void *waitpid_test(int arg1, void *arg2) {
    do_exit(arg1);
    return NULL;
}
```




## faber\_thread\_test()

- ➡ Let's look at the first set of tests in `faber_thread_test()`
- ▬ 2nd subtest

```
/**
 * Call do_waitpid with the process ID of the given
 * process. Print a debug message with the exiting
 * process's status.
 */
static void wait_for_proc(proc_t *p) {
    int rv;
    pid_t pid;
    char pname[PROC_NAME_LEN];

    strncpy(pname, p->p_comm, PROC_NAME_LEN);
    pname[PROC_NAME_LEN-1] = '\0';
    pid = do_waitpid(p->p_pid, 0, &rv);
    dbg(DBG_TEST, "%s (%d) exited: %d\n", pname, pid, rv);
}
```




## faber\_thread\_test ()

- ➡ Let's look at the first set of tests in `faber_thread_test ()`
- ▬ 3rd subtest

```
dbg (DBG_TEST, "kthread exit test\n");  
start_proc (&pt, "kthread exit test",  
            kthread_exit_test, 0);  
wait_for_proc (pt.p);
```

```
/*  
 * A thread function that returns NULL, silently  
 * invoking kthread_exit()  
 */  
void *kthread_exit_test (int arg1, void *arg2) {  
    return NULL;  
}
```






## faber\_thread\_test()

➡ Let's look at the first set of tests in `faber_thread_test()`  
 = 4th subtest

```
dbg(DBG_TEST, "many test\n");
for (i = 0; i < 10; i++)
    start_proc(NULL, "many test", waitpid_test, i);
wait_for_all();
dbg(DBG_TEST, " (C.1) done\n");
```

```
void *waitpid_test(int arg1, void *arg2) {
    do_exit(arg1);
    return NULL;
}
```


```
/* Repeatedly call wait_for_any() until it returns
 * -ECHILD */
static void wait_for_all() {
    while (wait_for_any() != -ECHILD)
        ;
}
```



## faber\_thread\_test()


- ➡ Let's look at the first set of tests in `faber_thread_test()`
- ▬ 4th subtest

```
/**
 * Call waitpid with a -1 pid and print a message
 * about any process that exits.
 * Returns the pid found, including -ECHILD when this
 * process has no children.
 */
static pid_t wait_for_any() {
    int rv;
    pid_t pid = do_waitpid(-1, 0, &rv);
    if (pid != -ECHILD)
        dbg(DBG_TEST, "child (%d) exited: %d\n", pid, rv);
    return pid;
}
```



# More Test Code


- ➡ We will not put all the code in "faber\_test.c" and "sunghan\_test.c" on these slides
  - in the discussion sections, we will discuss more of these test code
    - we will not cover all the test code
      - ◆ you need to learn how to read all these code and use them to figure out what you need to do
      - ◆ we will probably cover a *small number of test cases*
    - if you have a specific test you want the TA to talk about, please send him an e-mail before 5pm on Thursday
- ➡ **IMPORTANT:** at any line in `faber_thread_test()`, ask yourself
  - 1) where are all the threads/processes?
    - ◆ i.e., in which *queue* is a thread sleeping
  - 2) if a thread is not in the run queue, *who* is going to *unblock* it and *when/how* (by calling what function)?



# Kernel 2


Bill Cheng

*<http://merlot.usc.edu/william/usc/>*



# New Things in Kernel 2

- ➡ Polymorphism - VFS should be able to work with *any* AFS
  - achieved using *polymorphism*
    - the AFS for kernel 2 is `ramfs`
    - read the kernel code to see how it works
- ➡ Reference counting
  - whenever you *keep* a reference to an object, you increase the reference count of that object
    - so that the object won't disappear
  - whenever you remove a reference to an object, you decrement the reference count of that object
    - if it reached zero, it means that nothing in the system knows about that object and you should delete/free it
  - this is done for objects in memory
  - this is done for objects inside the file system (which you don't have to worry about)




# Kernel 2

% make nyi

|                      |                      |     |
|----------------------|----------------------|-----|
| main/kmain.c:218     | idleproc_run()       | VFS |
| main/kmain.c:223     | idleproc_run()       | VFS |
| fs/vnode.c:460       | special_file_read()  | VFS |
| fs/vnode.c:473       | special_file_write() | VFS |
| fs/vfs_syscall.c:68  | do_read()            | VFS |
| fs/vfs_syscall.c:83  | do_write()           | VFS |
| fs/vfs_syscall.c:97  | do_close()           | VFS |
| fs/vfs_syscall.c:120 | do_dup()             | VFS |
| fs/vfs_syscall.c:136 | do_dup2()            | VFS |
| fs/vfs_syscall.c:168 | do_mknod()           | VFS |
| fs/vfs_syscall.c:189 | do_mkdir()           | VFS |
| fs/vfs_syscall.c:214 | do_rmdir()           | VFS |
| fs/vfs_syscall.c:235 | do_unlink()          | VFS |
| fs/vfs_syscall.c:263 | do_link()            | VFS |
| fs/vfs_syscall.c:278 | do_rename()          | VFS |
| fs/vfs_syscall.c:298 | do_chdir()           | VFS |
| fs/vfs_syscall.c:320 | do_getdent()         | VFS |
| fs/vfs_syscall.c:337 | do_lseek()           | VFS |
| fs/vfs_syscall.c:357 | do_stat()            | VFS |
| fs/namev.c:45        | lookup()             | VFS |
| fs/namev.c:72        | dir_namev()          | VFS |
| fs/namev.c:90        | open_namev()         | VFS |
| fs/open.c:94         | do_open()            | VFS |

— make sure PROCS=1, DRIVERS=1, VFS=1 in Config.mk



# Kernel 2



## Create devices, set current working directory

main/kmain.c:218

idleproc\_run()

VFS

main/kmain.c:223

idleproc\_run()

VFS



## Reading from device and writing to device

fs/vnode.c:460

special\_file\_read()

VFS

fs/vnode.c:473

special\_file\_write()

VFS



## Pathname resolution functions

fs/namev.c:45

lookup()

VFS

fs/namev.c:72


dir\_namev()

VFS

fs/namev.c:90

open\_namev()

VFS



# Kernel 2

## ➡ System calls


```
fs/vfs_syscall.c:68
fs/vfs_syscall.c:83
fs/vfs_syscall.c:97
fs/vfs_syscall.c:120
fs/vfs_syscall.c:136
fs/vfs_syscall.c:168
fs/vfs_syscall.c:189
fs/vfs_syscall.c:214
fs/vfs_syscall.c:235
fs/vfs_syscall.c:263
fs/vfs_syscall.c:278
fs/vfs_syscall.c:298
fs/vfs_syscall.c:320
fs/vfs_syscall.c:337
fs/vfs_syscall.c:357
```

```
do_read()      VFS
do_write()     VFS
do_close()     VFS
do_dup()       VFS
do_dup2()      VFS
do_mknod()     VFS
do_mkdir()     VFS
do_rmdir()     VFS
do_unlink()    VFS
do_link()      VFS
do_rename()    VFS
do_chdir()     VFS
do_getdent()   VFS
do_lseek()     VFS
do_stat()      VFS
```

## ➡ Open

```
fs/open.c:94
```


```
do_open()      VFS
```






# Kernel 2: Where To Start

- ➡ **(Phase 1)** Create directories and devices in `idleproc_run()`
  - ▬ `do_mkdir()` and `do_mknod()`
    - and whatever you need to support these functions (including *pathname resolution* functions in "namev.c")
      - ◆ `open_namev()`, `dir_namev()`, `lookup()`
  - ▬ maybe get `do_stat()` to work to check your work
- ➡ **(Phase 2)** Get kshell to work again in `initproc_run()`
  - ▬ kernel 2 kshell is different from kernel 1 kshell
    - kernel 2 kshell uses VFS!
  - ▬ need to get a few more functions to work so that you can interact with kshell
    - `do_open()`, `do_write()`, `do_read()`,  
`special_file_write()`, `special_file_read()`
- ➡ **(Phase 3)** Create *kshell command* to run "vfstest"
  - ▬ pass all tests in "vfstest.c"
  - ▬ also need to pass tests in "faber\_fs\_test.c"




## vnode.c

```
/*  
 * If the file is a byte device then find the file's  
 * bytedev_t, and call read on it. Return what read  
 * returns.  
 *  
 * If the file is a block device then return -ENOTSUP  
 */  
static int  
special_file_read(vnode_t *file, off_t offset,  
                  void *buf, size_t count)  
{  
    NOT_YET_IMPLEMENTED("VFS: special_file_read");  
    return 0;  
}
```




## vnode.c

```
/*
 * If the file is a byte device find the file's
 * bytedev_t, and call its write. Return what write
 * returns.
 *
 * If the file is a block device then return -ENOTSUP.
 */
static int
special_file_write(vnode_t *file, off_t offset,
                  const void *buf, size_t count)
{
    NOT_YET_IMPLEMENTED("VFS: special_file_write");
    return 0;
}
```



## namev.c

```
/*  
 * This takes a base 'dir', a 'name', its 'len', and a  
 * result vnode.  
 *  
 * Most of the work should be done by the vnode's  
 * implementation specific lookup() function.  
 *  
 * If dir has no lookup(), return -ENOTDIR.  
 *  
 * Note: returns with the vnode refcount on *result  
 * incremented.  
 */  
int lookup(vnode_t *dir,  
           const char *name,  
           size_t len,  
           vnode_t **result);
```




# namev.c

```


/*
 * When successful this function returns data in the following
 * "out"-arguments:
 *   o res_vnode: the vnode of the parent directory of "name"
 *   o name: the 'basename' (the element of the pathname)
 *   o namelen: the length of the basename
 *
 * For example: dir_namev("/s5fs/bin/ls", &namelen, &name, NULL,
 * &res_vnode) would put 2 in namelen, "ls" in name, and a pointer to
 * the vnode corresponding to "/s5fs/bin" in res_vnode.
 *
 * The "base" argument defines where we start resolving the path from:
 * A base value of NULL means to use the process's current working
 * directory, curproc->p_cwd. If pathname[0] == '/', ignore base and
 * start with vfs_root_vn. dir_namev() should call lookup() to take
 * care of resolving each piece of the pathname.
 *
 * Note: A successful call to this causes vnode refcount on *res_vnode
 * to be incremented.
 */
int dir_namev(const char *pathname,
              size_t *namelen,
              const char **name,
              vnode_t *base,
              struct vnode **res_vnode);

```




## namev.c

```
/*  
 * This returns in res_vnode the vnode requested by the  
 * other parameters. It makes use of dir_namev and  
 * lookup to find the specified vnode (if it exists).  
 * flag is right out of the parameters to open(2); see  
 * <weenix/fcntl.h>. If the O_CREAT flag is specified,  
 * and the file does not exist call create() in the  
 * parent directory vnode.  
 *  
 * Note: Increments vnode refcount on *res_vnode.  
 */  
int open_namev(const char *pathname,  
               int flag,  
               vnode_t **res_vnode,  
               vnode_t *base);
```



## vfs\_syscall.c

```
int do_close(int fd);
int do_read(int fd, void *buf, size_t nbytes);
int do_write(int fd, const void *buf, size_t nbytes);
int do_dup(int fd);
int do_dup2(int ofd, int nfd);
int do_mknod(const char *path, int mode, unsigned devid);
int do_mkdir(const char *path);
int do_rmdir(const char *path);
int do_unlink(const char *path);
int do_link(const char *from, const char *to);
int do_rename(const char *oldname, const char *newname);
int do_chdir(const char *path);
int do_getdent(int fd, struct dirent *dirp);
int do_lseek(int fd, int offset, int whence);
int do_stat(const char *path, struct stat *uf);
```




# open.c

```

/*
 * There are a number of steps to opening a file:
 *
 * 1. Get the next empty file descriptor.
 *
 * 2. Call fget to get a fresh file_t.
 *
 * 3. Save the file_t in curproc's file descriptor table.
 *
 * 4. Set file_t->f_mode to OR of FMODE_(READ|WRITE|APPEND) based
 *    on oflags, which can be O_RDONLY, O_WRONLY or O_RDWR, possibly
 *    OR'd with O_APPEND or O_CREAT.
 *
 * 5. Use open_namev() to get the vnode for the file_t.
 *
 * 6. Fill in the fields of the file_t.
 *
 * 7. Return new fd.
 *
 * If anything goes wrong at any point (specifically if the call to
 * open_namev fails), be sure to remove the fd from curproc, fput the
 * file_t and return an error.
 *
 * Error cases you must handle for this function at the VFS level:
 *
 *   o EINVAL
 *   o EMFILE
 *   o ENOMEM
 *   o ENAMETOOLONG
 *   o ENOENT
 *   o EISDIR
 *   o ENXIO
 */
int do_open(const char *filename, int flags);

```






# Kernel 3


Bill Cheng

*<http://merlot.usc.edu/william/usc/>*



# New Things in Kernel 3


- ➡ **Polymorphism go recursive!**
  - fortunately, any recursion can be turned into a loop
    - so, think about the recursion as a loop
    - when you traverse a linked-list of shadow objects, you will eventually reach *bottom object* (and the recursion must stop)
- ➡ **Address space**
  - address space is implemented with *virtual memory maps*
    - review Ch 7 slides and *read kernel 3 FAQ*
- ➡ **Handle page faults**
  - keep Ch 7 slides in mind
- ➡ **System call and system call support**
  - `fork()` - where everything comes together
- ➡ **Debugging**
  - you may have to single-step machine instructions to find bugs



# Kernel 3

% make nyi


|                    |                          |    |
|--------------------|--------------------------|----|
| mm/pframe.c:359    | pframe_get()             | VM |
| mm/pframe.c:379    | pframe_pin()             | VM |
| mm/pframe.c:395    | pframe_unpin()           | VM |
| api/syscall.c:77   | sys_read()               | VM |
| api/syscall.c:87   | sys_write()              | VM |
| api/syscall.c:103  | sys_getdents()           | VM |
| api/access.c:144   | addr_perm()              | VM |
| api/access.c:160   | range_perm()             | VM |
| proc/fork.c:77     | do_fork()                | VM |
| vm/pagefault.c:72  | handle_pagefault()       | VM |
| vm/shadow.c:73     | shadow_init()            | VM |
| vm/shadow.c:85     | shadow_create()          | VM |
| vm/shadow.c:97     | shadow_ref()             | VM |
| vm/shadow.c:111    | shadow_put()             | VM |
| vm/shadow.c:126    | shadow_lookuppage()      | VM |
| vm/shadow.c:144    | shadow_fillpage()        | VM |
| vm/shadow.c:153    | shadow_dirtypage()       | VM |
| vm/shadow.c:160    | shadow_cleanpage()       | VM |
| proc/kthread.c:161 | kthread_clone()          | VM |
| fs/vnode.c:487     | special_file_mmap()      | VM |
| fs/vnode.c:499     | special_file_fillpage()  | VM |
| fs/vnode.c:511     | special_file_dirtypage() | VM |
| fs/vnode.c:523     | special_file_cleanpage() | VM |



# Kernel 3

% make nyi

|                |                        |    |
|----------------|------------------------|----|
| vm/vmmap.c:129 | vmmap_create()         | VM |
| vm/vmmap.c:138 | vmmap_destroy()        | VM |
| vm/vmmap.c:148 | vmmap_insert()         | VM |
| vm/vmmap.c:161 | vmmap_find_range()     | VM |
| vm/vmmap.c:171 | vmmap_lookup()         | VM |
| vm/vmmap.c:182 | vmmap_clone()          | VM |
| vm/vmmap.c:215 | vmmap_map()            | VM |
| vm/vmmap.c:251 | vmmap_remove()         | VM |
| vm/vmmap.c:262 | vmmap_is_range_empty() | VM |
| vm/vmmap.c:277 | vmmap_read()           | VM |
| vm/vmmap.c:292 | vmmap_write()          | VM |
| vm/brk.c:76    | do_brk()               | VM |
| vm/anon.c:60   | anon_init()            | VM |
| vm/anon.c:72   | anon_create()          | VM |
| vm/anon.c:84   | anon_ref()             | VM |
| vm/anon.c:98   | anon_put()             | VM |
| vm/anon.c:106  | anon_lookuppage()      | VM |
| vm/anon.c:115  | anon_fillpage()        | VM |
| vm/anon.c:122  | anon_dirtypage()       | VM |
| vm/anon.c:129  | anon_cleanpage()       | VM |
| vm/mmap.c:55   | do_mmap()              | VM |
| vm/mmap.c:70   | do_munmap()            | VM |



# /usr/bin/hello

## ➡ Page frame management

|                 |                |    |
|-----------------|----------------|----|
| mm/pframe.c:359 | pframe_get()   | VM |
| mm/pframe.c:379 | pframe_pin()   | VM |
| mm/pframe.c:395 | pframe_unpin() | VM |


## ➡ User address space implementation

|                |                        |    |
|----------------|------------------------|----|
| vm/vmmap.c:129 | vmmap_create()         | VM |
| vm/vmmap.c:138 | vmmap_destroy()        | VM |
| vm/vmmap.c:148 | vmmap_insert()         | VM |
| vm/vmmap.c:161 | vmmap_find_range()     | VM |
| vm/vmmap.c:171 | vmmap_lookup()         | VM |
| vm/vmmap.c:182 | vmmap_clone()          | VM |
| vm/vmmap.c:215 | vmmap_map()            | VM |
| vm/vmmap.c:251 | vmmap_remove()         | VM |
| vm/vmmap.c:262 | vmmap_is_range_empty() | VM |
| vm/vmmap.c:277 | vmmap_read()           | VM |
| vm/vmmap.c:292 | vmmap_write()          | VM |

- some of them are needed by the loader in "api/elf32.c" since the loader creates an address space for a user program

## ➡ Handle page faults

|                   |                    |    |
|-------------------|--------------------|----|
| vm/pagefault.c:72 | handle_pagefault() | VM |
|-------------------|--------------------|----|



# /usr/bin/hello

## ➡ Basic system calls and system call support


|                   |                |    |
|-------------------|----------------|----|
| api/syscall.c:77  | sys_read()     | VM |
| api/syscall.c:87  | sys_write()    | VM |
| api/syscall.c:103 | sys_getdents() | VM |
| api/access.c:144  | addr_perm()    | VM |
| api/access.c:160  | range_perm()   | VM |

## ➡ Memory management (mmobj)

|               |                   |    |
|---------------|-------------------|----|
| vm/anon.c:60  | anon_init()       | VM |
| vm/anon.c:72  | anon_create()     | VM |
| vm/anon.c:84  | anon_ref()        | VM |
| vm/anon.c:98  | anon_put()        | VM |
| vm/anon.c:106 | anon_lookuppage() | VM |
| vm/anon.c:115 | anon_fillpage()   | VM |
| vm/anon.c:122 | anon_dirtypage()  | VM |
| vm/anon.c:129 | anon_cleanpage()  | VM |

### — recall that there are 3 types of mmobj

- one lives inside the vnode
- anonymous object
- shadow object



# /usr/bin/hello

## ➡ Basic system calls and system call support

|                   |                |
|-------------------|----------------|
| api/syscall.c:77  | sys_read()     |
| api/syscall.c:87  | sys_write()    |
| api/syscall.c:103 | sys_getdents() |

— this may not be the complete list (I was told that these are what you need)

|                  |              |    |
|------------------|--------------|----|
| api/access.c:144 | addr_perm()  | VM |
| api/access.c:160 | range_perm() | VM |

## ➡ Memory management (mmobj)

|               |                   |    |
|---------------|-------------------|----|
| vm/anon.c:60  | anon_init()       | VM |
| vm/anon.c:72  | anon_create()     | VM |
| vm/anon.c:84  | anon_ref()        | VM |
| vm/anon.c:98  | anon_put()        | VM |
| vm/anon.c:106 | anon_lookuppage() | VM |
| vm/anon.c:115 | anon_fillpage()   | VM |
| vm/anon.c:122 | anon_dirty_page() | VM |
| vm/anon.c:129 | anon_cleanpage()  | VM |


"hello"



— recall that there are 3 types of mmobj

- one lives inside the vnode
- anonymous object
- shadow object

➡ To get "hello" to work, you need to get all of the above to work (except for sys\_read() and sys\_getdents())



# Beyond /usr/bin/hello

## ➡ More memory management (mmobj)


|                 |                     |    |
|-----------------|---------------------|----|
| vm/shadow.c:73  | shadow_init()       | VM |
| vm/shadow.c:85  | shadow_create()     | VM |
| vm/shadow.c:97  | shadow_ref()        | VM |
| vm/shadow.c:111 | shadow_put()        | VM |
| vm/shadow.c:126 | shadow_lookuppage() | VM |
| vm/shadow.c:144 | shadow_fillpage()   | VM |
| vm/shadow.c:153 | shadow_dirty_page() | VM |
| vm/shadow.c:160 | shadow_cleanpage()  | VM |

### — shadow object is complicated

- only really need this when you get to run "fork-and-wait"

## ➡ More system calls and system call support

|                    |                           |    |
|--------------------|---------------------------|----|
| proc/fork.c:77     | do_fork()                 | VM |
| proc/kthread.c:161 | kthread_clone()           | VM |
| vm/mmap.c:55       | do_mmap()                 | VM |
| vm/mmap.c:70       | do_munmap()               | VM |
| vm/brk.c:76        | do_brk()                  | VM |
| fs/vnode.c:487     | special_file_mmap()       | VM |
| fs/vnode.c:499     | special_file_fillpage()   | VM |
| fs/vnode.c:511     | special_file_dirty_page() | VM |
| fs/vnode.c:523     | special_file_cleanpage()  | VM |





# Beyond /usr/bin/hello

## ➡ More memory management (mmobj)

|                 |                     |    |
|-----------------|---------------------|----|
| vm/shadow.c:73  | shadow_init()       | VM |
| vm/shadow.c:85  | shadow_create()     | VM |
| vm/shadow.c:97  | shadow_ref()        | VM |
| vm/shadow.c:111 | shadow_put()        | VM |
| vm/shadow.c:126 | shadow_lookuppage() | VM |
| vm/shadow.c:144 | shadow_fillpage()   | VM |
| vm/shadow.c:153 | shadow_dirtypage()  | VM |
| vm/shadow.c:160 | shadow_cleanpage()  | VM |

— shadow object is complicated

- only really need this when you get to run "**fork-and-wait**"


## ➡ More system calls and system call support

|                    |                 |    |
|--------------------|-----------------|----|
| proc/fork.c:77     | do_fork()       | VM |
| proc/kthread.c:161 | kthread_clone() | VM |

---

|              |             |    |
|--------------|-------------|----|
| vm/mmap.c:55 | do_mmap()   | VM |
| vm/mmap.c:70 | do_munmap() | VM |
| vm/brk.c:76  | do_brk()    | VM |

|                |                          |    |
|----------------|--------------------------|----|
| fs/vnode.c:487 | special_file_mmap()      | VM |
| fs/vnode.c:499 | special_file_fillpage()  | VM |
| fs/vnode.c:511 | special_file_dirtypage() | VM |
| fs/vnode.c:523 | special_file_cleanpage() | VM |



# Beyond /usr/bin/hello

## ➡ More memory management (mmobj)

|                 |                     |    |
|-----------------|---------------------|----|
| vm/shadow.c:73  | shadow_init()       | VM |
| vm/shadow.c:85  | shadow_create()     | VM |
| vm/shadow.c:97  | shadow_ref()        | VM |
| vm/shadow.c:111 | shadow_put()        | VM |
| vm/shadow.c:126 | shadow_lookuppage() | VM |
| vm/shadow.c:144 | shadow_fillpage()   | VM |
| vm/shadow.c:153 | shadow_dirty_page() | VM |
| vm/shadow.c:160 | shadow_cleanpage()  | VM |


### — shadow object is complicated

- only really need this when you get to run "fork-and-wait"

## ➡ More system calls and system call support

|                    |                           |    |
|--------------------|---------------------------|----|
| proc/fork.c:77     | do_fork()                 | VM |
| proc/kthread.c:161 | kthread_clone()           | VM |
| vm/mmap.c:55       | do_mmap()                 | VM |
| vm/mmap.c:70       | do_munmap()               | VM |
| vm/brk.c:76        | do_brk()                  | VM |
| fs/vnode.c:487     | special_file_mmap()       | VM |
| fs/vnode.c:499     | special_file_fillpage()   | VM |
| fs/vnode.c:511     | special_file_dirty_page() | VM |
| fs/vnode.c:523     | special_file_cleanpage()  | VM |

**"/sbin/init"**




# Strategy

➡ You must implement page frame management first


|                 |                |    |
|-----------------|----------------|----|
| mm/pframe.c:359 | pframe_get()   | VM |
| mm/pframe.c:379 | pframe_pin()   | VM |
| mm/pframe.c:395 | pframe_unpin() | VM |

- set S5FS=1 but keep VM=0 in Config.mk (even though this is kernel 3)
  - VM=1 means that you are running user-space programs
- the re-run vfstest and make sure everything works
  - s5fs has a real disk and ramfs didn't
- don't bother with tests in "faber\_fs\_test.c" because they were designed for kernel 2 only
- you should get this done by *end of week 12*
  - if code is mostly working but cannot past all vfstest, split up your team
    - ◆ one to debug code in "pframe.c", the rest move forward




# Strategy

- ➡ In the *first week* of kernel 3, you should get the *first user-space program* (i.e., "hello") to run to run
- ➡ in `initproc_run()`, run `"/usr/bin/hello"` using `kernel_execve()`
    - the kernel's `INIT` process becomes the user-space "hello" process
    - you have to be able to *load* the program into memory
    - you have to be able to *build and manipulate* the *address space*
    - you have to be able to *handle page faults*
    - you have to get some *system call* (such as `write()`) to work
  - ➡ that's a lot of stuff to get working just to run "hello"
    - please do *not* even attempt to run `"/sbin/init"`
  - ➡ would be good to get this done by the *end of week 1 of kernel 3*



# Strategy

- ➡ In the *2nd week* of kernel 3, you should get all the user-space programs in section (B) of the grading guidelines running from `initproc_run()`
- ➡ the last one is `"/usr/bin/fork_and_wait"`
    - you need to get `fork()` to work
  - ➡ by the end of the weekend, you should get `"/sbin/init"` to run from `initproc_run()`
    - the kernel's INIT process becomes the user-space INIT process
      - ◆ the user-space INIT process spawn child processes to run user-space shell (`"/bin/sh"`)
      - ◆ from the user-space shell (`"/bin/sh"`), you should be able to re-run all the section (B) user programs
  - ➡ would be good to get this done by the *end of weekend of the 2nd week of kernel 3*
- ➡ In the *last week* of kernel 3, get everything else to work




# pframe.c

```

/*
 * Find and return the pframe representing the page identified by the
 * object and page number. If the page is already resident in memory,
 * then we return the existing page. Otherwise, we allocate a new page
 * and fill it (in which case this routine may block). Before allocating
 * the new pframe, we check to see if we need to call pageoutd and wake
 * it up if necessary.
 *
 * If the page is found (resident) but busy, then we will wait for it to
 * become unbusy and then try again (since it may have been freed after
 * that). Thus, as long as this routine returns successfully, the returned
 * page will be a non-busy page that will be guaranteed to remain resident
 * until the calling context blocks without first pinning the page.
 *
 * This routine may block at the mmobj operation level.
 *
 * @param o the parent object of the page
 * @param pagenum the page number of this page in the object
 * @param result used to return the pframe (NULL if there's an error)
 * @return 0 on success, < 0 on failure.
 */
int
pframe_get(struct mmobj *o, uint32_t pagenum, pframe_t **result)
{
    NOT_YET_IMPLEMENTED("VM: pframe_get");
    return 0;
}

```




# pframe.c

```


/*
 * Increases the pin count on this page. Pages with a pin count > 0 will
 * not be paged out by pageoutd, so this ensures that the page will remain
 * resident until the pin count is decreased.
 *
 * If the pframe has not yet been pinned, remove this pframe's list link
 * from the allocated list and add it to the pinned list. Be sure to
 * decrement nallocated and increment npinned.
 *
 * In either case, increment the pf_pincount.
 *
 * @param pf the page to pin
 */
int
pframe_pin(pframe_t *pf)
{
    NOT_YET_IMPLEMENTED("VM: pframe_pin");
    return 0;
}

```



# pframe.c

```
/*
 * Decreases the pin count on a page. If the pin count reaches zero,
 * then the page could be paged out any time after the calling context
 * blocks.
 *
 * If the pin count reaches zero, move the pframe's list link from the
 * pinned list to the allocated list. Be sure to correctly update
 * npinned and nallocated
 *
 * @param pf a pinned page (a page with a positive pin count)
 */
int
pframe_unpin(pframe_t *pf)
{
    NOT_YET_IMPLEMENTED("VM: pframe_unpin");
    return 0;
}
```






# More On Kernel 3

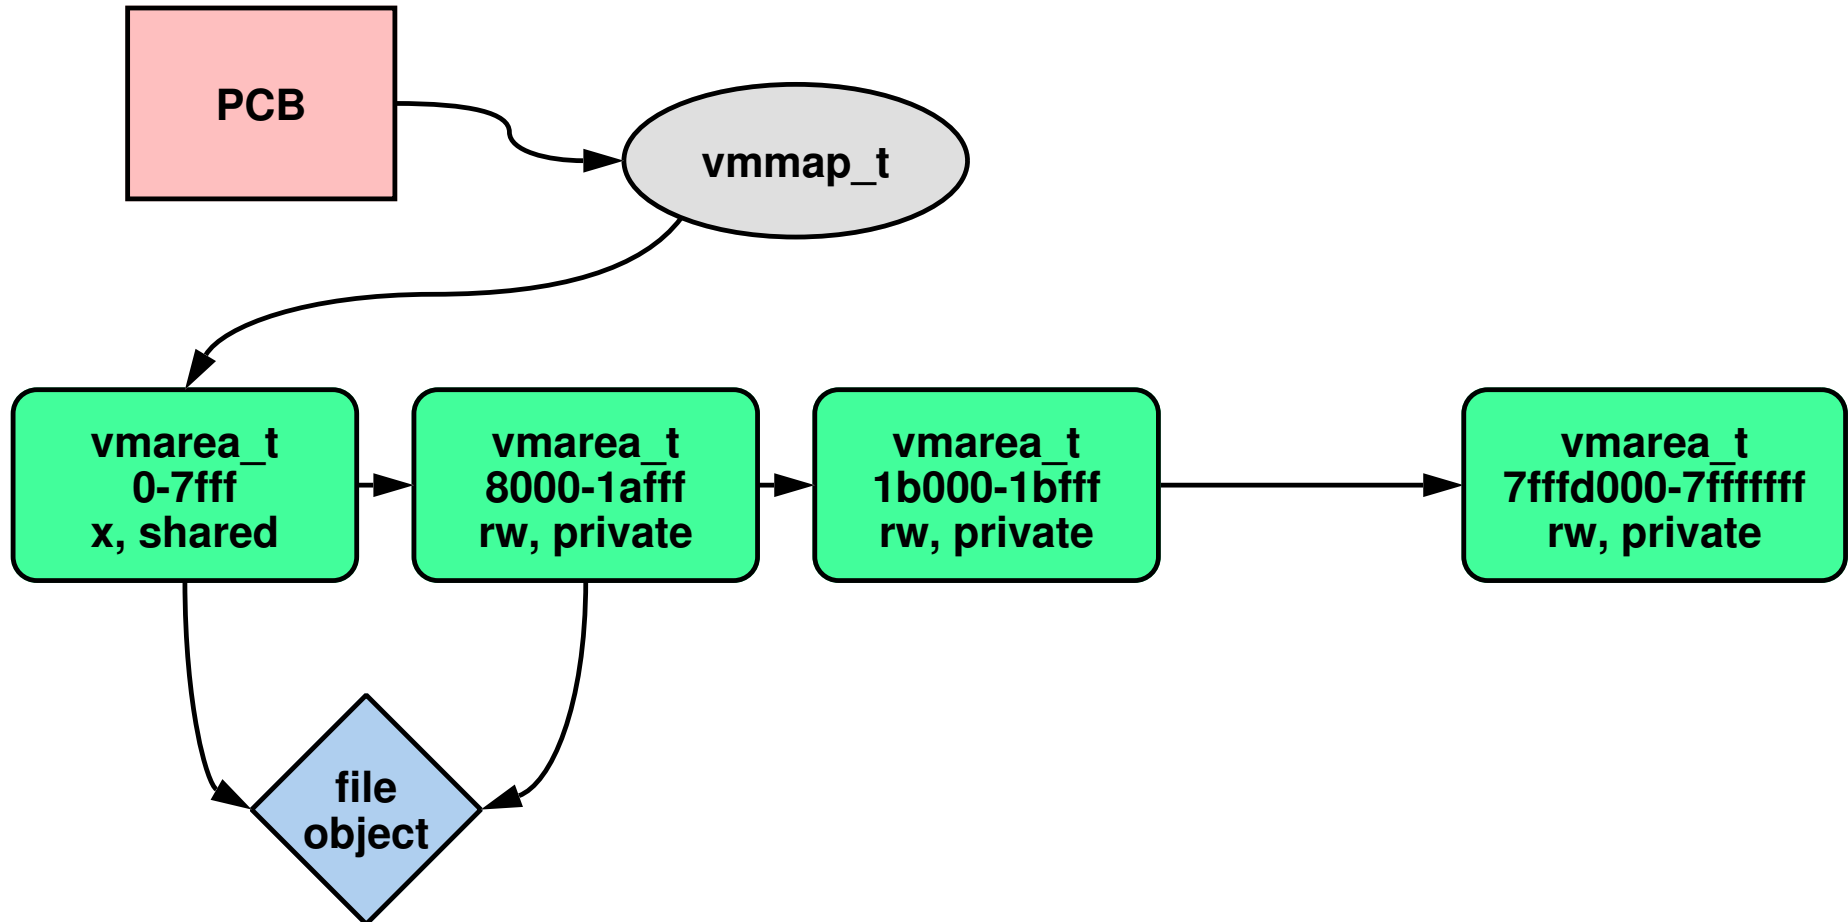
Bill Cheng

*<http://merlot.usc.edu/william/usc/>*




# Address Space Implementation

- ➡ **Address Space** is implemented using **Virtual Memory Map** (`vmmmap`)
- in lectures, sometimes I used the term "memory map"

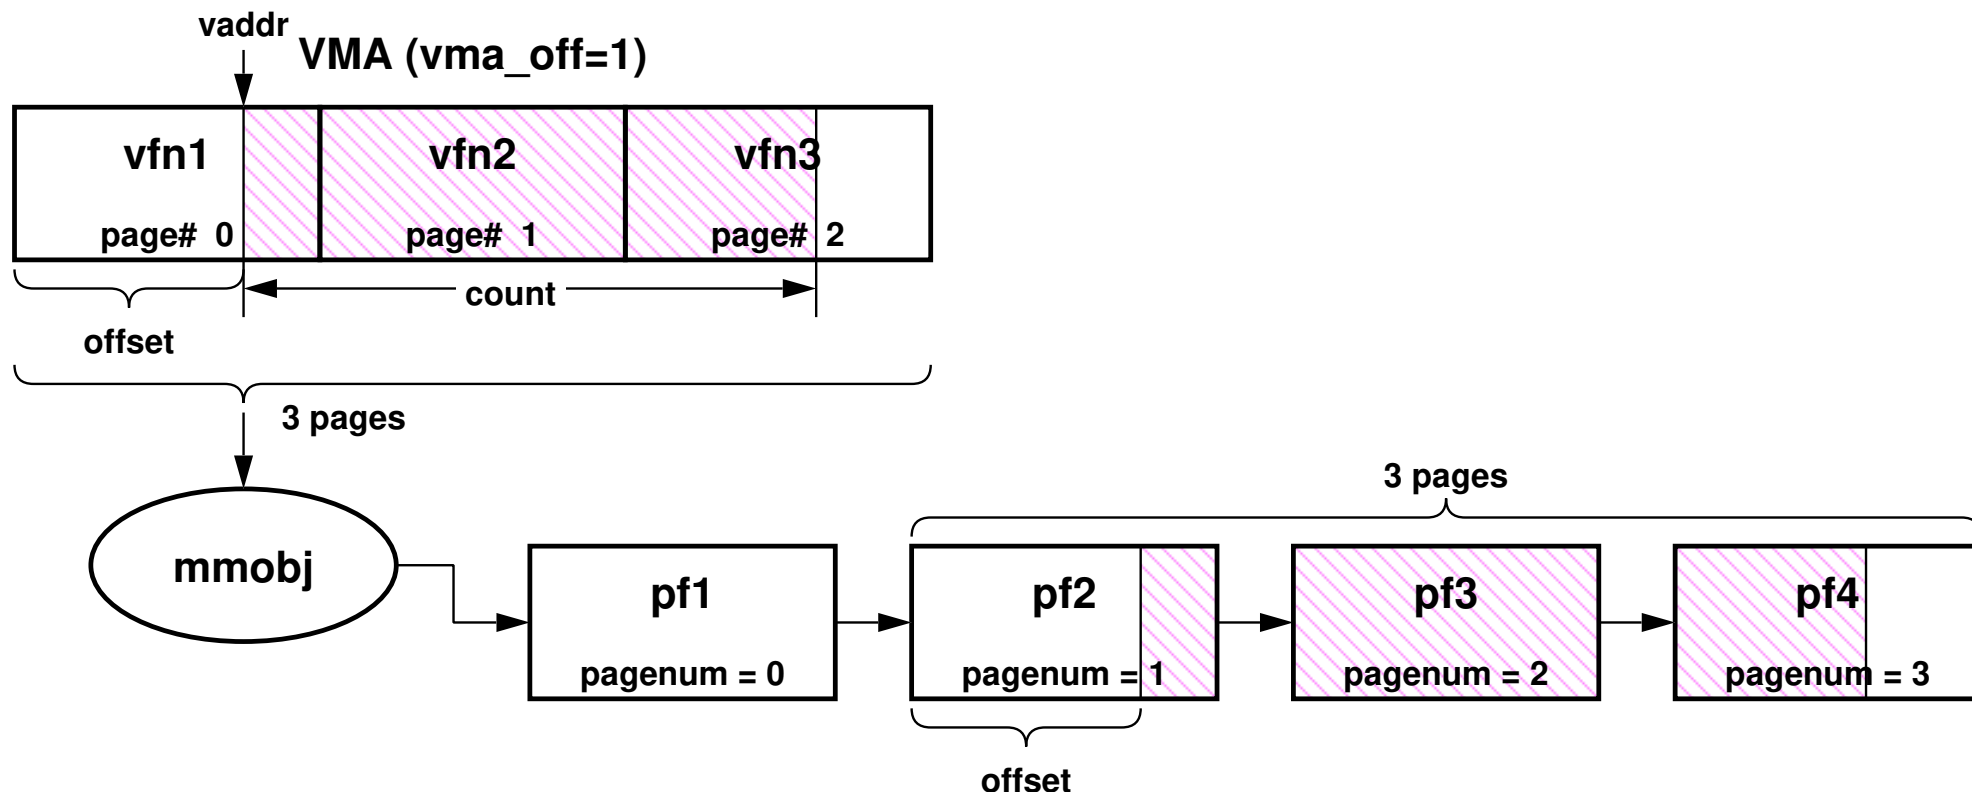


- we will use `vma` or `vmarea` to refer to `vmarea_t`
- the values in `vmas` above are not what's in weenix




# mmobj and pagenum


➡ This is a very important picture in the kernel FAQ to understand




- read the kernel FAQ about "pagenum" and the next few FAQ items that follow it
- these are crucial in understanding how to build an address space for a user process



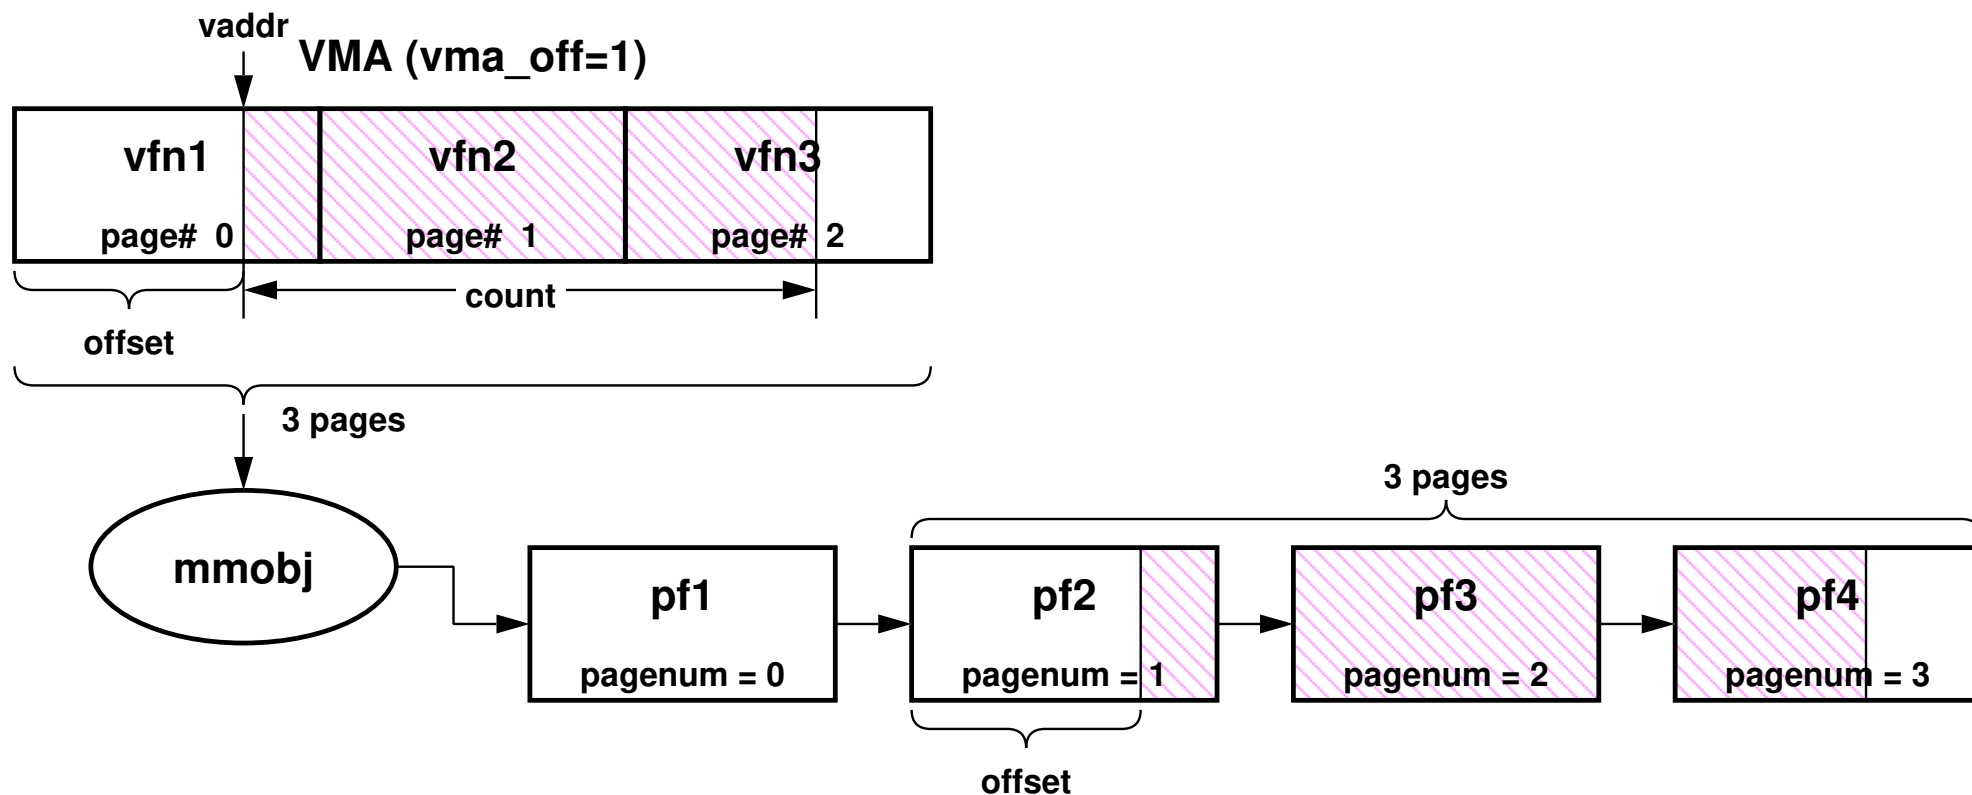
# mmobj and pagenum



- 1) The address space is made up of a list of *non-overlapping* vmareas
- 2) Each vmarea is a memory segment (contiguous virtual memory locations)
  - in the above, it's shown as  $[vaddr, vaddr + count)$ 
    - typically  $vaddr$  (first virtual address of a memory segment) is *page aligned*




# mmobj and pagenum

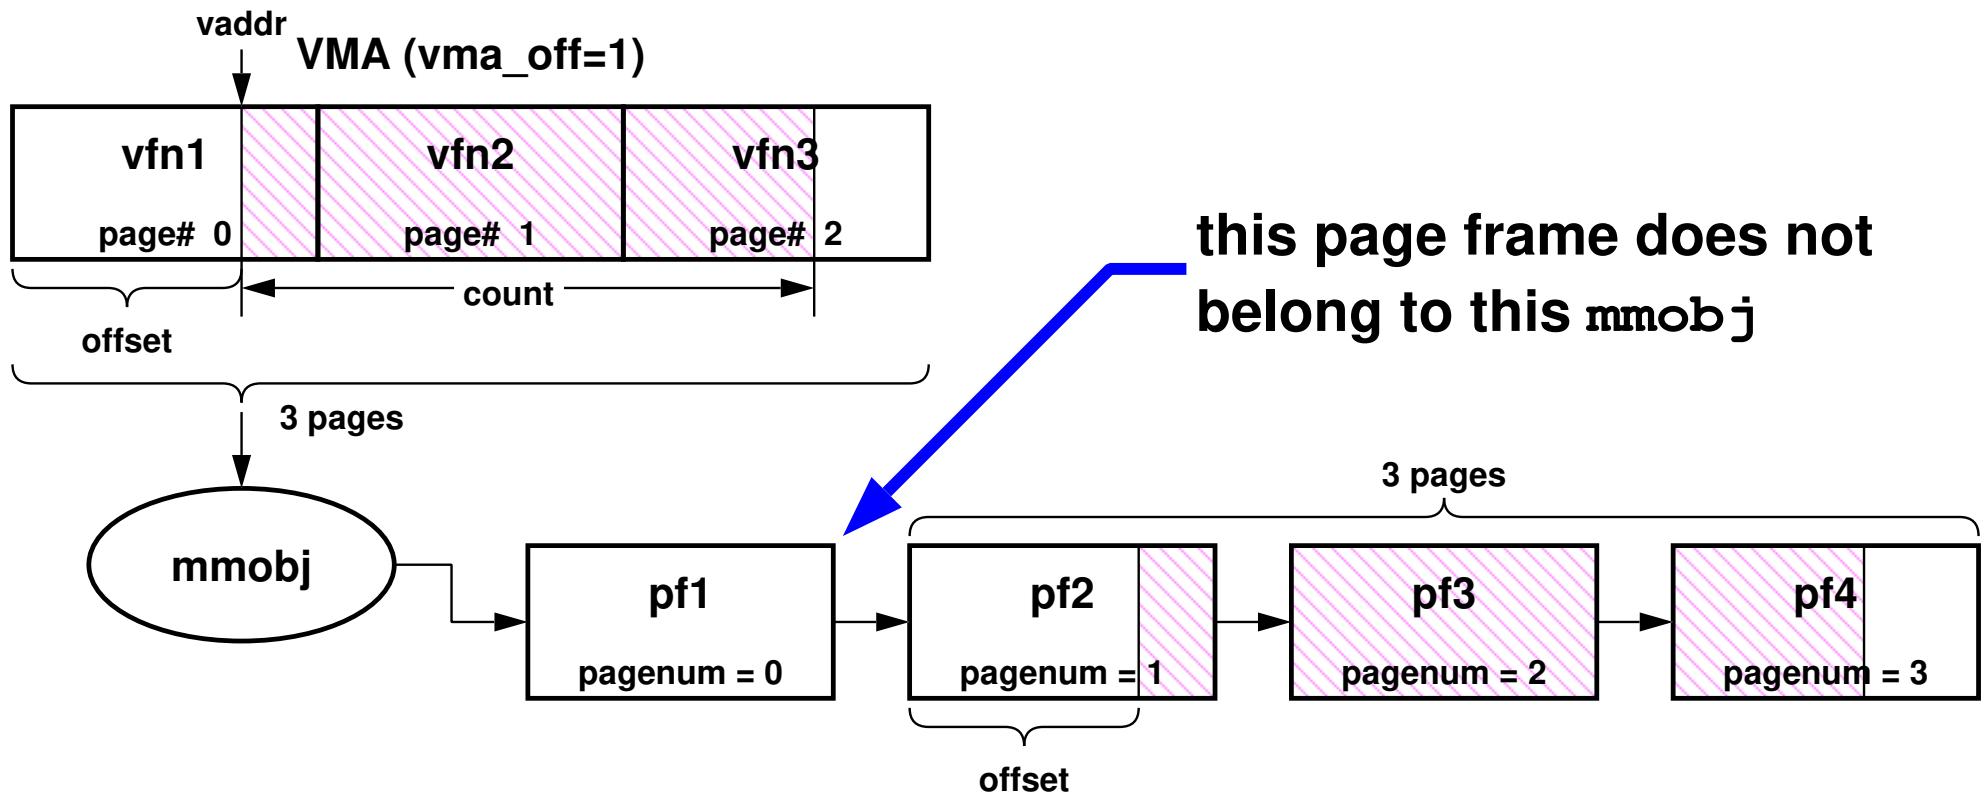


3) The kernel manages VM in "pages" (not "bytes"), it allocates enough pages so that a memory segment can fit inside

- ➡ above, it takes 3 virtual pages to cover this memory segment
- ➡ each virtual page need to be mapped into a physical page
  - vfn = vpn = virtual page number (20-bits long)
  - you need to get used to printing these numbers in *hex*




# mmobj and pagenum

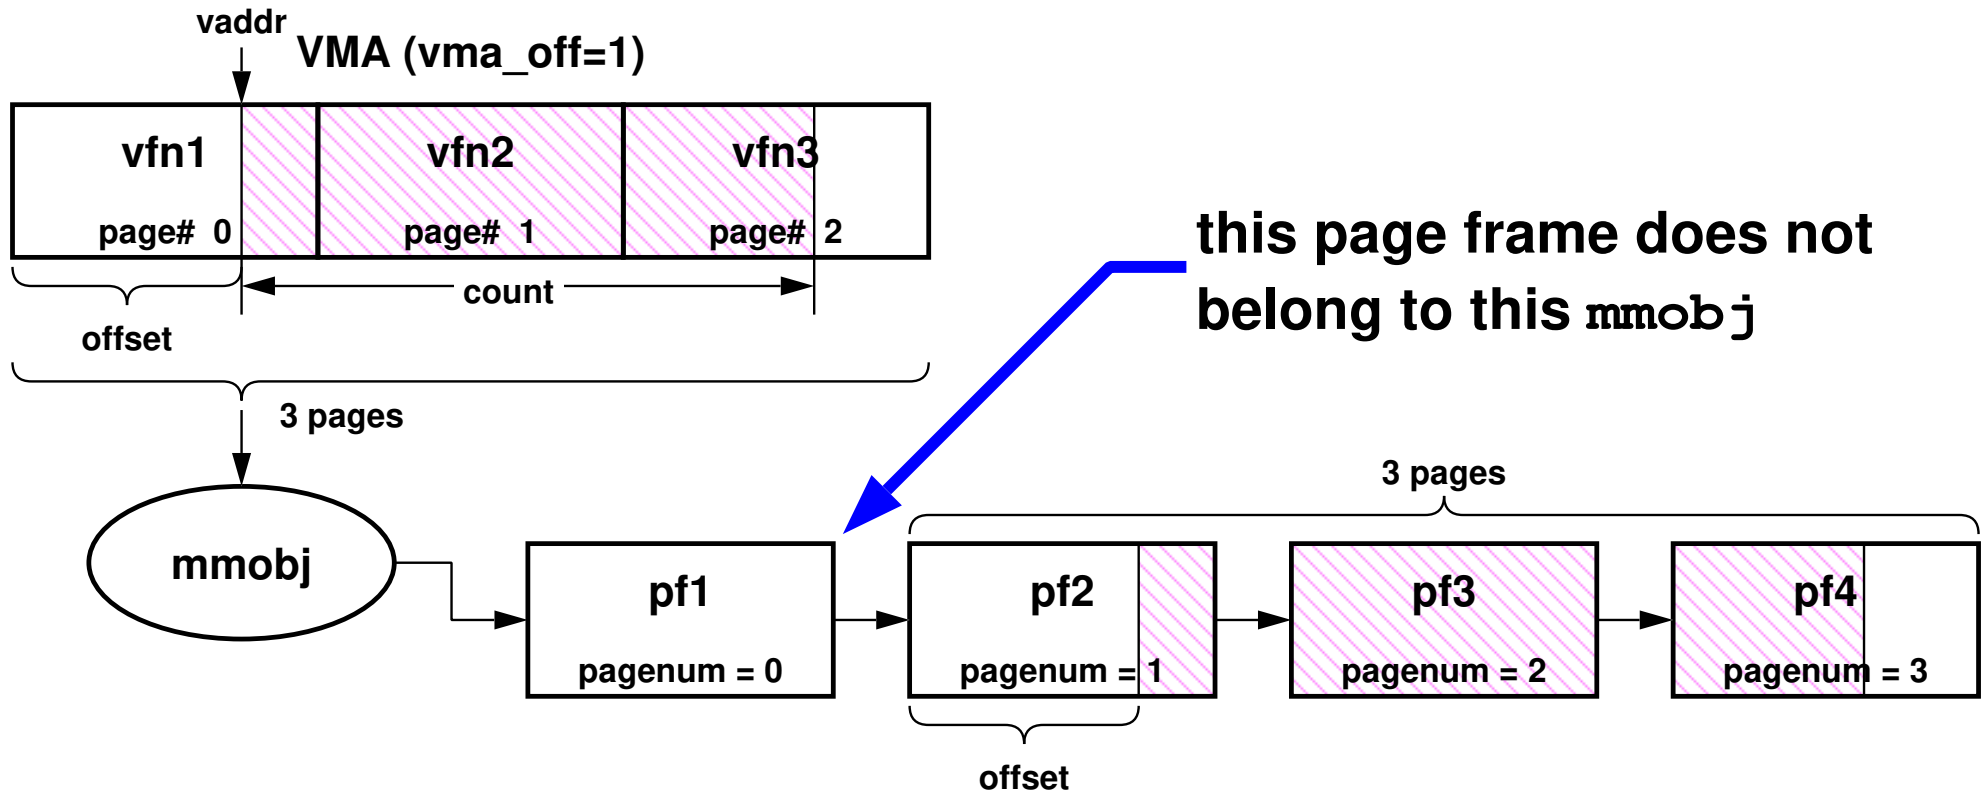


## 4) Page frames are managed by mmobjs

- there is a physical page (hidden) inside each page frame
- if an mmobj manages N page frames, their pagenums are  $[vma\_off, vma\_off+N)$  for that mmobj
- *not really implemented as a linked list* as shown above
  - a *hash table* is used




# mmobj and pagenum




5) You can create multiple memory segments by mapping different pages of a file into your address space


- ▢ conceptually, a file is divided into pages
- ▢ **vma\_off** in a memory segment gives you the starting page number of the file



# mmobj and pagenum

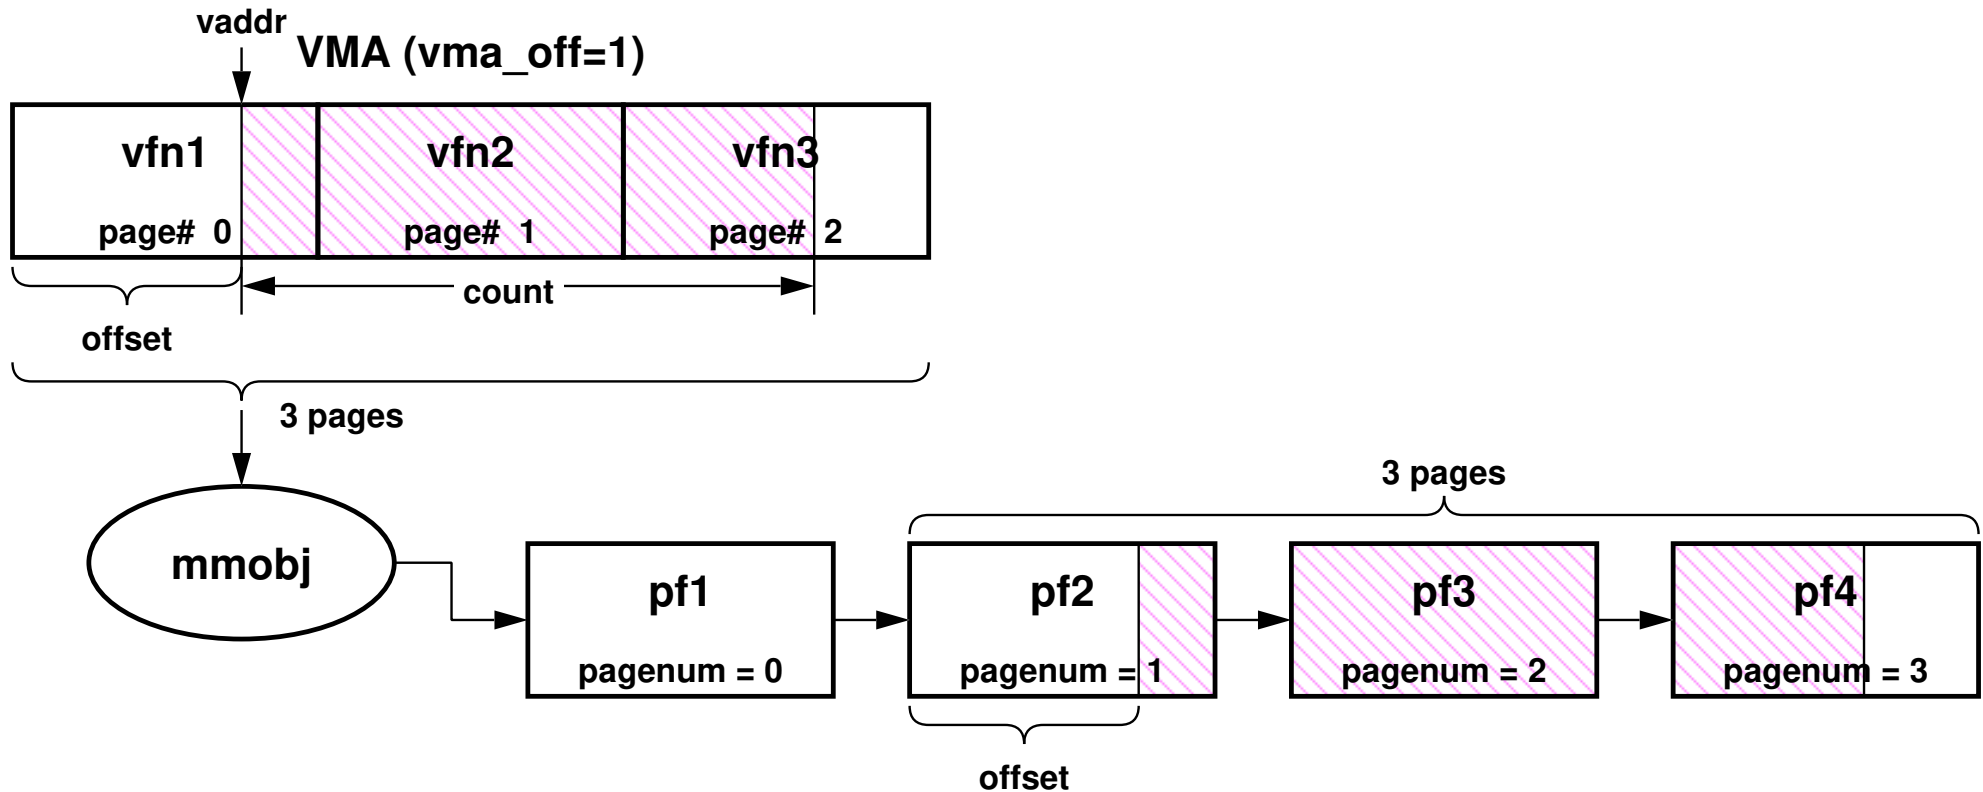


- 6) A page frame may not be present for an `mmobj`
- ⇒ since we are doing *demand paging*, in the beginning, none of the page frames are present for an `mmobj`
  - ⇒ as page frames are brought in, they become "*resident*"
    - i.e., they are *cached* inside the corresponding `mmobj`
    - if modified, the page frame becomes "*dirty*"






# mmobj and pagenum

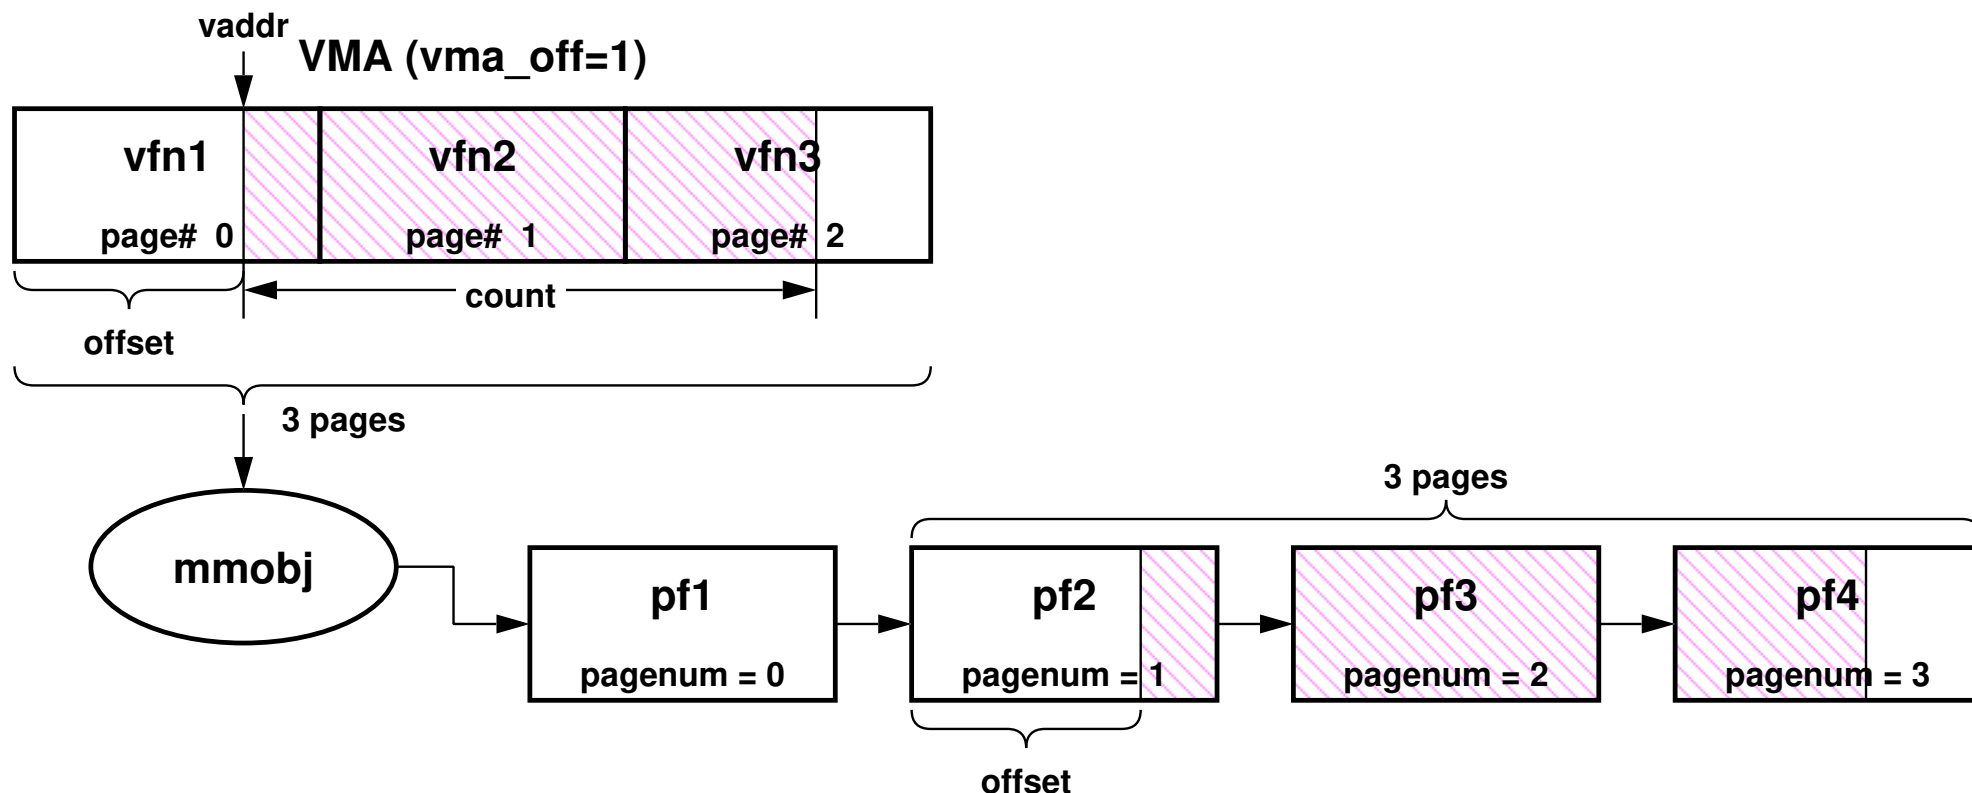


7) A page frame is identified by the `mmobj` that manages it and the `pagenum` of the page frame


- ⇒ in the kernel FAQ, it uses the notation **(o,n)** where **o** is an `mmobj` and **n** is a `pagenum`
- ⇒ a **linked list** of `mmobjs` must be used to make **copy-on-write** work correctly with `fork()`



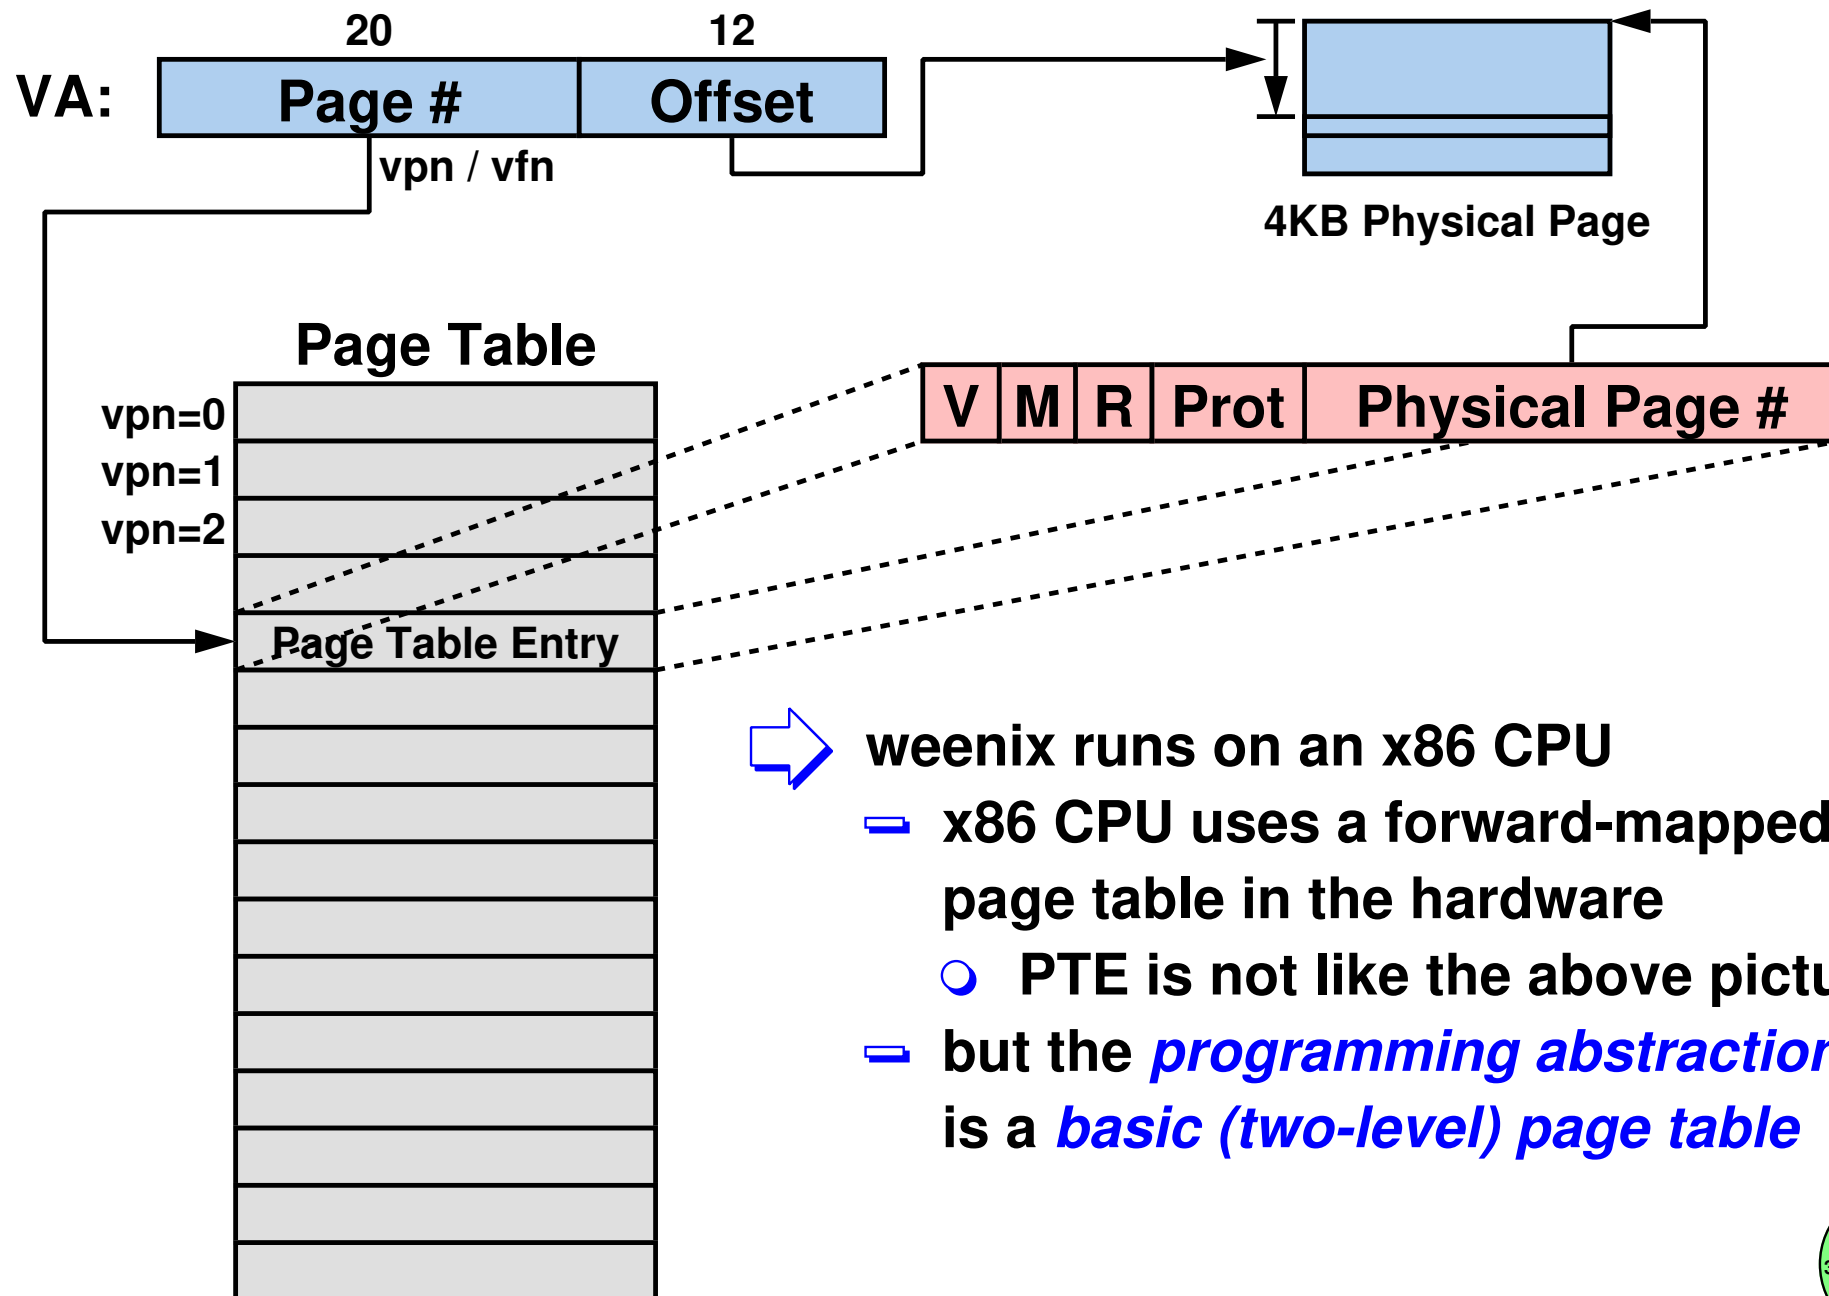
# mmobj and pagenum



- 8) Page table maps a **virtual page number** to a **physical page number**
- ⇒ in the above picture, vfn1/vfn2/vfn3 needs to be mapped to the physical page that lives inside pf2/pf3/pf4, respectively
  - ⇒ pf\_addr of a page frame contains a **kernel virtual address** for the corresponding "physical page"
  - pt\_virt\_to\_phys() converts it to physical address




# Forward-Mapped (Multilevel) Page Table



weenix runs on an x86 CPU

- x86 CPU uses a forward-mapped page table in the hardware
- PTE is not like the above picture
- but the *programming abstraction* is a *basic (two-level) page table*



# Very Useful gdb Commands


## ➡ Address space

```
kernel info vmmap_mapping_info curproc->p_vmmmap
```

- ➡ the *loader* builds address space for a user-space program
- ➡ after your user-space program is "loaded", what does the address space look like?
- ➡ when you get your first legitimate page fault in `handle_pagefault()`, look at your address space
  - if it's wrong, is there point proceeding?!
  - ask your classmates in the class Google Group if they are seeing the same thing
    - ◆ don't just ask others to share, that's not nice!

## ➡ Page table (not that useful)

```
kernel info pt_mapping_info curproc->p_pagedir
```



# Read Some Code

## ➡ User

— "hello.c"

```
int main(int argc, char **argv)
{
    open("/dev/tty0", O_RDONLY, 0);
    open("/dev/tty0", O_WRONLY, 0);
    write(1, "Hello, world!\n", 14);
    return 0;
}
```

## ➡ Kernel

- "elf32.c" - the loader
- "access.c" - need to understand `copy_from_user()` and `copy_to_user()`
- "exec.c" - how to go into user space
- "pagefault.c" - handle page fault
- "vn\_mmobj\_ops.c" - code for the `mmobj` inside a `vnode`
- "syscall.c" - to see how to implement `sys_write()`, look at how other `sys_*` functions are implemented

